

Implementing Domain-Driven Design with Java

Premanand Chandrasekaran<[@premanandc](https://github.com/premanandc)>
and Karthik Krishnan<[@karthik](https://github.com/karthik-krishnan)>

Version 1.0, 2020-11-02

Table of Contents

Preface.....	1
Who This Book Is For	1
What You Will Learn	1
Acknowledgements	1
Part 1: Foundations	2
1. The Rationale for Domain-Driven Design	3
1.1. Introduction	3
1.2. Why do software projects fail?	3
1.2.1. Inaccurate requirements	4
1.2.2. Too much architecture	4
1.2.3. Too little architecture	5
1.2.4. Excessive incidental complexity	5
1.2.5. Uncontrolled technical debt	6
1.2.6. Ignoring Non-Functional Requirements (NFRs)	6
1.2.7. Where To From Here?.....	7
1.3. Modern Systems and Dealing with Complexity	7
1.4. What is Domain-Driven Design?.....	9
1.4.1. What is a Domain?.....	9
1.4.2. What is a Subdomain?.....	10
1.4.3. Types of Subdomains.....	11
1.4.4. Domain Experts	12
1.4.5. Promoting a Shared Understanding	12
1.4.6. Evolving a Domain Model and a Solution	12
1.4.7. The Essence of DDD	13
1.5. Why is DDD Relevant? Why Now?.....	14
1.5.1. Rise of Open Source	14
1.5.2. Advances in Technology	15
1.5.3. Rise of Distributed Computing	16
1.6. Summary	16
1.7. Questions	17
1.8. Further Reading	17
1.9. Answers	18
2. The Mechanics of Domain-Driven Design (30 Pages)	19
2.1. Understanding the Problem	19
2.1.1. Problem Space vs. Solution Space	19
2.1.2. Dealing with Ambiguity	19
2.2. Arriving at a Shared Understanding	19
2.3. Breaking Down the Problem	19

2.3.1. What is a Domain?	19
2.3.2. What is a Sub-Domain?	19
2.3.3. The Core Sub-Domain	19
2.4. Modeling a Solution	19
2.4.1. What is a Model?	19
2.4.2. Context Maps	19
2.4.3. Bounded Contexts	19
2.5. Implementing the Solution	19
3. Where and How Does DDD Fit? (15 pages)	20
3.1. Architecture Styles	20
3.2. Layered Architecture	20
3.3. Onion Architecture	20
3.4. Hexagonal Architecture	20
3.5. Service Oriented Architecture	20
3.6. Microservice Architecture	20
3.7. Event-Driven Architecture (EDA)	20
3.8. Command Query Responsibility Segregation (CQRS)	20
3.9. Serverless Architecture	20
Part 2: Implementing DDD in the Real World	21
4. Domain Analysis and Modeling	22
4.1. Introduction	22
4.2. Technical requirements	22
4.3. Introducing the LC application	22
4.3.1. What is a Letter of Credit (LC)	23
4.3.2. The LC issuance application	23
4.4. Enhancing shared understanding	23
4.5. Domain storytelling	24
4.5.1. Introducing Domain Storytelling	24
4.5.2. Using DST for the LC application	26
4.6. EventStorming	30
4.6.1. Introducing EventStorming	30
4.6.2. Using eventStorming for the LC issuance application	31
4.7. Summary	34
4.8. Questions	34
4.9. Further Reading	35
4.10. Answers	35
5. Implementing Domain Logic	36
5.1. Technical Requirements	36
5.2. Command Query Responsibility Segregation (CQRS)	36
5.2.1. Recap: What is CQRS	37
5.2.2. Why CQRS?	37

5.3. Bootstrapping the application	38
5.4. Identifying Commands	39
5.4.1. Identifying Aggregates	40
5.4.2. Test-driving the system	42
5.4.3. Implementing the command	43
5.4.4. Implementing the event	44
5.4.5. Designing the aggregate	45
5.5. Persisting aggregates	47
5.5.1. State stored aggregates	47
5.5.2. Event sourced aggregates	48
5.5.3. Which persistence mechanism should we choose?	51
5.6. Policy enforcements (validations)	52
5.6.1. Structural validations	52
5.6.2. Business rule enforcements	53
5.7. Summary	61
5.8. Questions	61
5.9. Further reading	61
5.10. Answers	62
6. Implementing the User Interface—Task-based (20 pages)	63
6.1. Technical Requirements	63
6.2. Bootstrapping the UI	63
6.3. API Styles	66
6.3.1. CRUD-based APIs	67
6.3.2. Task-based APIs	69
6.3.3. Impact of API style on user experiences	70
6.3.4. Task-based or CRUD-based?	70
6.4. Design pattern considerations	70
6.4.1. Usability	70
6.4.2. Testability	70
6.5. Implementing the save LC functionality	70
6.6. Invoking backend services	70
6.7. Summary	70
6.8. Questions	70
6.9. Further reading	70
7. Implementing Queries and Projections (10 pages)	71
7.1. Consuming events	71
7.2. Persisting Distinct Query Models	71
7.3. Exposing a REST-based API for Queries	71
7.4. Creating Additional Views	71
8. Long-Running User Flows (10 pages)	72
8.1. Implementing Sagas	72

8.2. Taking Care of Deadlines	72
8.3. Distributed Exception Handling	72
8.4. Keeping Track of the Overall Flow	72
8.5. Deciding between Orchestration and Choreography	72
9. Integrating with External Systems (15 pages)	73
9.1. Technical Requirements	73
9.2. Implementing consumer-driven contracts	73
9.3. Exposing a REST-based API	73
9.4. Exposing an events-based API	73
9.5. Implementing an Anti-Corruption Layer	73
9.6. Legacy Application Migration Patterns	73
Part 3: Advanced Patterns	74
10. Distributing into Microservices (15 pages)	75
10.1. Right Sizing Components	75
10.2. Maintaining Autonomy	75
10.3. Understanding the Costs of Distribution	75
10.4. Testing the Overall System	75
11. Non-Functional Requirements (25 pages)	76
11.1. Dealing With Eventual Consistency	76
11.2. Scaling the Event Store with Snapshots	76
11.3. Event Versioning and Upcasting	76
11.4. Monitoring, Metrics and Tracing	76
11.5. Enhancing Performance	76
12. Migrating to Serverless (15 pages)	77
12.1. Serverless Primer	77
12.2. Services as Functions	77
12.3. Serverless Persistence	77
12.4. Next Steps	77

Preface

Domain-Driven Design makes available a set of techniques and patterns that non-technical experts, architects and developers to work together and decompose complex systems into well-factored, collaborating, loosely coupled subsystems. Write more here... TODO.

Who This Book Is For

Developers working with Domain-Driven Design will be able to put their knowledge to work with this practical guide to create elegant software designs that are pleasant to work with and easy to work with and reason about. The book provides a hands-on approach to implementation and associated methodologies that will have you up-and-running, and productive in no time.

What You Will Learn

By the end of this book, you will be able to architect, design and implement robust, modern and loosely coupled distributed architectures employing domain-driven design.

Acknowledgements

TODO

Part 1: Foundations

While the IT industry prides itself on being at the very bleeding edge of technology, it also oversees a relatively high proportion of projects that fail outright or do not meet their originally intended goals for one reason or another. In Part 1, we will look at reasons for software projects not achieving their intended objectives and how practising Domain-Driven Design (DDD) can significantly help improve the odds of achieving success. We will also do a quick tour of the main concepts that Eric Evans elaborated in his seminal book by the same name and examine why/how it is extremely relevant in today's distributed systems age.

Chapter 1. The Rationale for Domain-Driven Design

The being cannot be termed rational or virtuous, who obeys any authority, but that of reason.

— Mary Wollstonecraft

1.1. Introduction

According to the Project Management Institute's (PMI) *Pulse of the Profession* report published in February 2020, only 77% of all projects meet their intended goals — and even this is true only in the most mature organizations. For less mature organizations, this number falls to just 56% i.e. approximately one in every two projects does not meet its intended goals. Furthermore, approximately one in every five projects is declared an outright failure. At the same time, we also seem to be embarking on our most ambitious and complex projects.

In this chapter, we will examine the main causes for project failure and look at how applying domain-driven design provides a set of guidelines and techniques to improve the odds of success in our favor. While Eric Evans wrote his classic book on the subject way back in 2003, we look at why that work is still extremely relevant in today's times.

1.2. Why do software projects fail?

Failure is simply the opportunity to begin again, this time more intelligently.

— Henry Ford

According to the [project success report](#) published in the Project Management Journal of the PMI, the following six factors need to be true for a project to be deemed successful:

Project Success Factors

Category	Criterion	Description
Project	Time	It meets the desired time schedules
	Cost	Its cost does not exceed budget
	Performance	It works as intended
Client	Use	Its intended clients use it
	Satisfaction	Its intended clients are happy
	Effectiveness	Its intended clients derive direct benefits through its implementation

With all of these criteria being applied to assess project success, a large percentage of projects fail for one reason or another. Let's examine some of the top reasons in more detail:

1.2.1. Inaccurate requirements

PMI's *Pulse of the Profession* report from 2017 highlights a very starking fact—a vast majority of projects fail due to inaccurate or misinterpreted requirements. It follows that it is impossible to build something that clients can use, are happy with and makes them more effective at their jobs if the wrong thing gets built—even much less for the project to be built on time, and under budget.

IT teams, especially in large organizations are staffed with mono-skilled roles such as UX designer, developer, tester, architect, business analyst, project manager, product owner, business sponsor, etc. In a lot of cases, these people are parts of distinct organization units/departments—each with its own set of priorities and motivations. To make matters even worse, the geographical separation between these people only keeps increasing. The need to keep costs down and the current COVID-19 ecosystem does not help matters either.

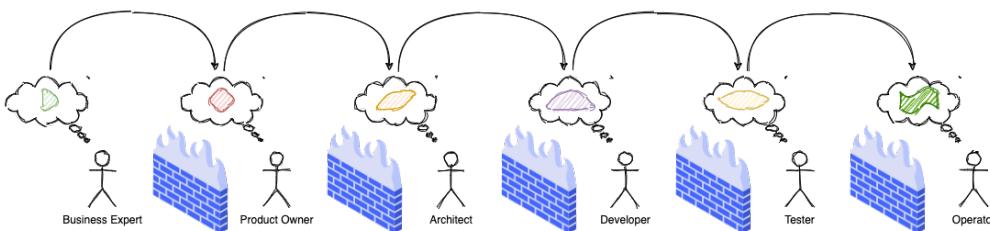


Figure 1- 1. Silo mentality and the loss of information fidelity

All this results in a loss in fidelity of information at every stage in the *assembly line*, which then results in misconceptions, inaccuracies, delays and eventually failure!

1.2.2. Too much architecture

Writing complex software is quite a task. One cannot just hope to sit down and start typing code—although that approach might work in some trivial cases. Before translating business ideas into working software, a thorough understanding of the problem at hand is necessary. For example, it is not possible (or at least extremely hard) to build credit card software without understanding how credit cards work in the first place. To communicate one's understanding of a problem, it is not uncommon to create software models of the problem, before writing code. This model or collection of models represents the understanding of the problem and the architecture of the solution.

Efforts to create a perfect model of the problem—one that is accurate in a very broad context, are not dissimilar to the proverbial holy grail quest. Those accountable to produce the architecture can get stuck in [analysis paralysis](#) and/or [big design up front](#), producing artifacts that are one or more of too high level, wishful, gold plated, buzzword-driven, disconnected from the real world—while not solving any real business problems. This kind of *lock-in* can be especially detrimental during the early phases of the project when knowledge levels of team members are still up and coming. Needless to say, projects adopting such approaches find it hard to meet with success consistently.



For a more comprehensive list of [modeling anti-patterns](#), refer to Scott W. Ambler's website (<http://agilemodeling.com>) and book dedicated to the subject.

1.2.3. Too little architecture

Agile software delivery methods manifested themselves in the late 90s, early 2000s in response to heavyweight processes collectively known as *waterfall*. These processes seemed to favor [big design up front](#) and abstract ivory tower thinking based on wishful, ideal world scenarios. This was based on the premise that thinking things out well in advance ends up saving serious development headaches later on as the project progresses.

In contrast, agile methods seem to favor a much more nimble and iterative approach to software development with a high focus on working software over other artifacts such as documentation. Most teams these days claim to practice some form of iterative software development. However, this obsession to claim conformance to a specific family of [agile methodologies](#) as opposed to the underlying principles, a lot of teams misconstrue having just enough architecture with having no perceptible architecture. This results in a situation where adding new features or enhancing existing ones takes a lot longer than what it previously used to—which then accelerates the devolution of the solution to become the dreaded [big ball of mud](#).

1.2.4. Excessive incidental complexity

Mike Cohn popularized the notion of the [test pyramid](#) where he talks about how a large number of unit tests should form the foundation of a sound testing strategy—with numbers decreasing significantly as one moves up the pyramid. The rationale here is that as one moves up the pyramid, the cost of upkeep goes up copiously while speed of execution slows down manifold. In reality though, a lot of teams seem to adopt a strategy that is the exact opposite of this—known as the testing ice cream cone as depicted below:



Figure 1-2. Testing Strategy: Expectation vs. Reality

The testing ice cream cone is a classic case of what Fred Brooks calls incidental complexity in his seminal paper titled [No Silver Bullet—Essence and Accident in Software Engineering](#). All software has some amount of [essential complexity](#) that is inherent to the problem being solved. This is especially true when creating solutions for non-trivial problems. However, incidental or accidental complexity is not directly attributable to the problem itself—but is caused by limitations of the people involved, their skill levels, the tools and/or abstractions being used. Not keeping tabs on incidental complexity causes teams to veer away from focusing on the real problems, solving which provide the most value. It naturally follows that such teams minimize their odds of success appreciably.

1.2.5. Uncontrolled technical debt

Financial debt is the act of borrowing money from an outside party to quickly finance the operations of a business—with the promise to repay the principal plus the agreed upon rate of interest in a timely manner. Under the right circumstances, this can accelerate the growth of a business considerably while allowing the owner to retain ownership, reduced taxes and lower interest rates. On the other hand, the inability to pay back this debt on time can adversely affect credit rating, result in higher interest rates, cash flow difficulties, and other restrictions.

Technical debt is what results when development teams take arguably sub-optimal actions to expedite the delivery of a set of features or projects. For a period of time, just like borrowed money allows you to do things sooner than you could otherwise, technical debt can result in short term speed. In the long term, however, software teams will have to dedicate a lot more time and effort towards simply managing complexity as opposed to thinking about producing architecturally sound solutions. This can result in a vicious negative cycle as illustrated in the diagram below:



Figure 1-3. Technical Debt—Implications

In a recent [McKinsey survey](#) sent out to CIOs, around 60% reported that the amount of tech debt increased over the past three years. At the same time, over 90% of CIOs allocated less than a fifth of their tech budget towards paying it off. Martin Fowler [explores](#) the deep correlation between high software quality (or the lack thereof) and the ability to enhance software predictably. While carrying a certain amount of tech debt is inevitable and part of doing business, not having a plan to systematically pay off this debt can have significantly detrimental effects on team productivity and ability to deliver value.

1.2.6. Ignoring Non-Functional Requirements (NFRs)

Stakeholders often want software teams to spend a majority (if not all) of their time working on features that provide enhanced functionality. This is understandable given that such features provide the highest ROI. These features are called functional requirements.

Non-functional requirements, on the other hand, are those aspects of the system that do not affect

functionality directly, but have a profound effect on the efficacy of those using these and maintaining these systems. There are many kinds of NFRs. A partial list of common NFRs is depicted below:



Figure 1- 4. Non-Functional Requirements

Very rarely do users explicitly request non-functional requirements, but almost always expect these features to be part of any system they use. Oftentimes, systems may continue to function without NFRs being met, but not without having an adverse impact on the *quality* of the user experience. For example, the home page of a web site that loads in under 1 second under low load and takes upwards of 30 seconds under higher loads may not be usable during those times of stress. Needless to say, not treating non-functional requirements with the same amount of rigor as explicit, value-adding functional features, can lead to unusable systems—and subsequently failure.

1.2.7. Where To From Here?

In this section we examined some common reasons that cause software projects to fail. In the upcoming section, we will look at characteristics of modern systems and look at more effective ways to deal with software complexity. In upcoming chapters, we will look at how applying domain-driven design helps mitigate these causes of failure.

1.3. Modern Systems and Dealing with Complexity

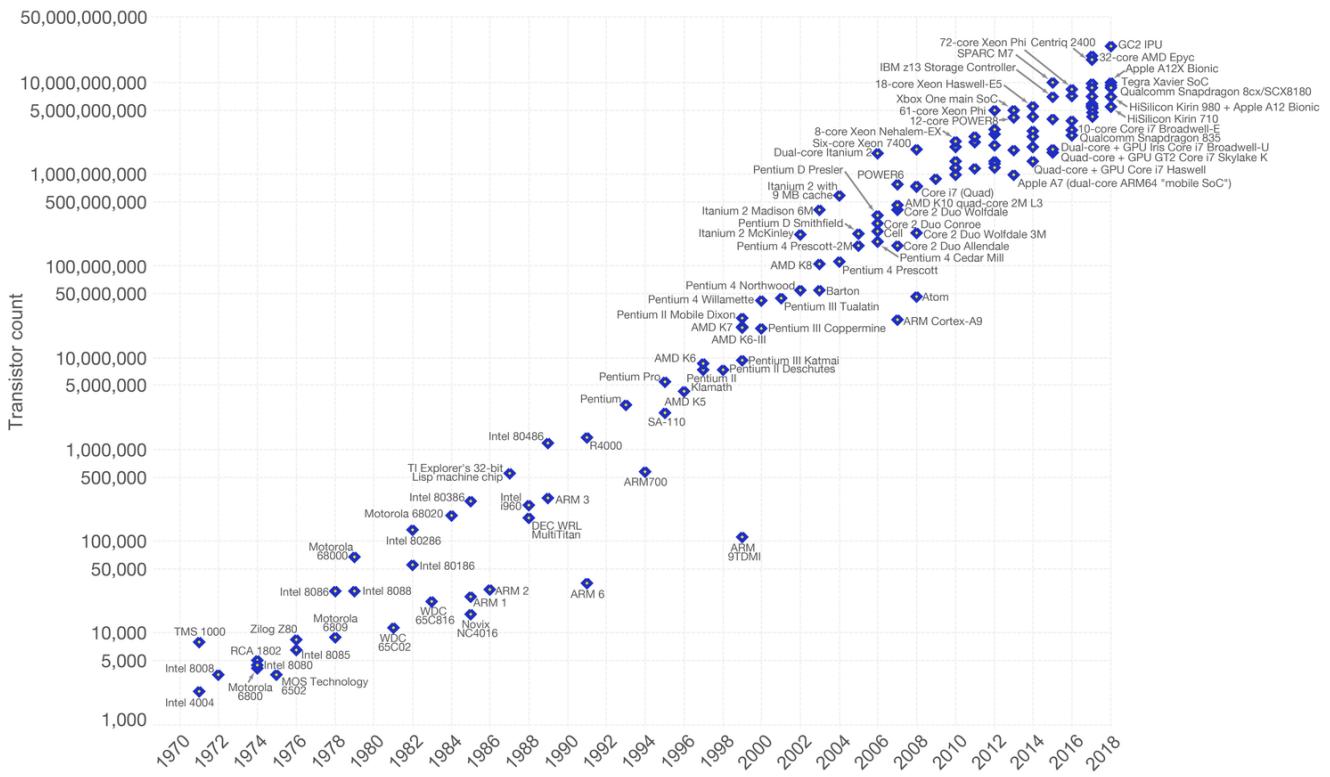
We can not solve our problems with the same level of thinking that created them.

— Albert Einstein

We find ourselves in the midst of the fourth industrial revolution where the world is becoming more and more digital—with technology being a significant driver of value for businesses. Exponential advances in computing technology as illustrated by Moore's Law below,

Moore's Law – The number of transistors on integrated circuit chips (1971-2018)

Moore's law describes the empirical regularity that the number of transistors on integrated circuits doubles approximately every two years. This advancement is important as other aspects of technological progress – such as processing speed or the price of electronic products – are linked to Moore's law.



Data source: Wikipedia (https://en.wikipedia.org/wiki/Transistor_count)
The data is available at [Gartner](https://www.semiconductorindustry.org/semiconductor-industry-statistics/). The following

The data visualization is available at OurWorldInData.org. There you find more visualizations and research on this topic.

Licensed under CC-BY-SA by the author Max Roser.

Figure 1-5. Moore's Law

along with the rise of the internet as illustrated below,

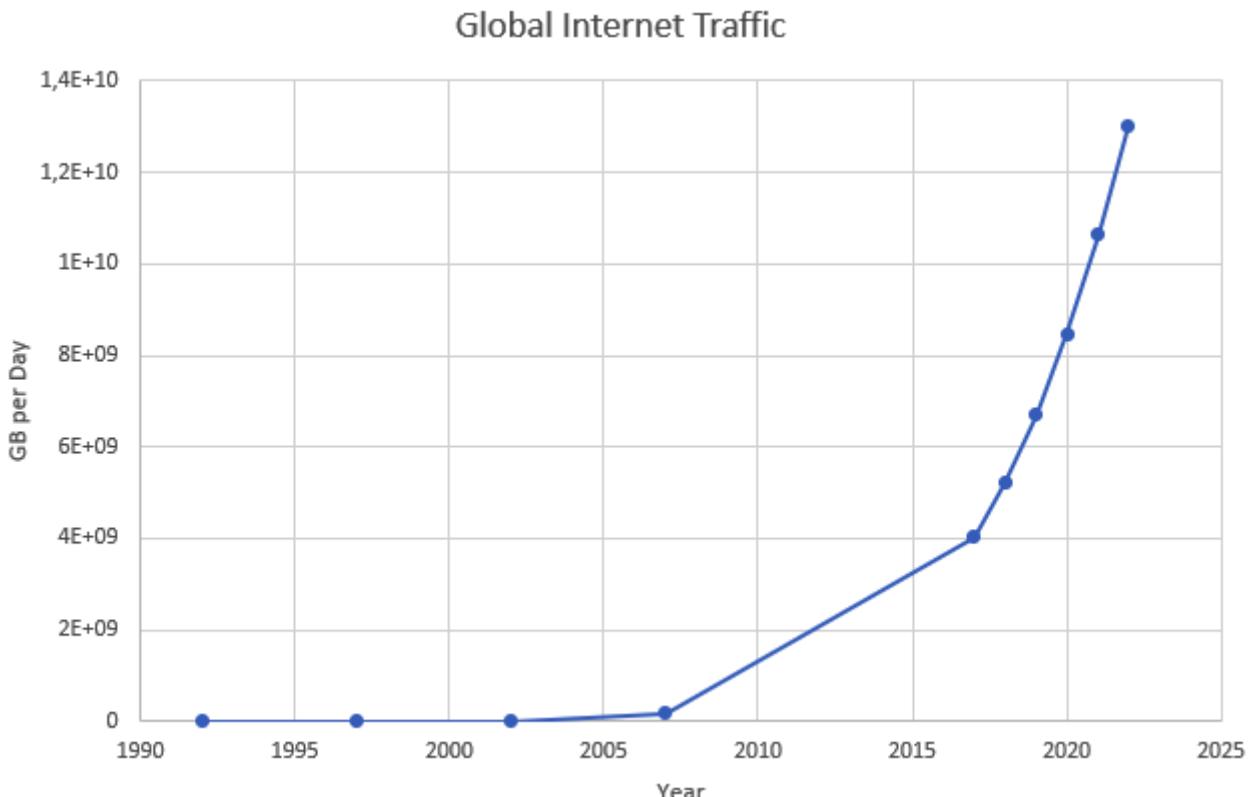


Figure 1- 6. Global Internet Traffic

has meant that companies are being required to modernize their software systems much more rapidly than they ever have. Along with all this, the onset of commodity computing services such as the public cloud has led to a move away from expensive centralized computing systems to more distributed computing ecosystems. As we attempt building our most complex solutions, monoliths are being replaced by an environment of distributed, collaborating microservices. Modern philosophies and practices such as automated testing, architecture fitness functions, continuous integration, continuous delivery, devops, security automation, infrastructure as code, to name a few, are disrupting the way we deliver software solutions.

As we enter an age of encountering our most complex business problems, we need to embrace new ways of thinking, a development philosophy and an arsenal of techniques to iteratively evolve mature software solutions that will stand the test of time. We need better ways of communicating, analyzing problems, arriving at a collective understanding, creating and modeling abstractions, and then implementing, enhancing the solution.

Domain-driven design promises to provide answers on how to do this in a systematic manner. In the upcoming section, and indeed the rest of this book, we will examine what DDD is and why it is indispensable when working to provide solutions for non-trivial problems in today's world of massively distributed teams and applications.

1.4. What is Domain-Driven Design?

Life is really simple, but we insist on making it complicated.

— Confucius

In the previous section, we saw how a myriad of reasons coupled with system complexity get in the way of software project success. The idea of domain-driven design, originally conceived by Eric Evans in his 2003 book, is an approach to software development that focuses on expressing software solutions in the form of a model that closely embodies the core of the problem being solved. It provides a set of principles and systematic techniques to analyze, architect and implement software solutions in a manner that enhances chances of success.

While Evans' work was indeed seminal, ground-breaking, and way ahead of its time, over the years, practical application has continued to remain a challenge. In this section, we will look at some of the foundational terms and concepts behind domain-driven design. Elaboration and practical application of these concepts will happen in upcoming chapters of this book.

To understand DDD, first and foremost, we need to understand what we mean by the first "D"—**domain**.

1.4.1. What is a Domain?

The foundational concept when working with domain-driven design is the notion of a domain. But what exactly is a domain? The word **domain**, which has its [origins](#) in the 1600s to the Old French word *domaine* (power), Latin word *dominium* (property, right of ownership) is a rather confusing word. Depending on who, when, where and how it is used, it can mean different things:

Noun [edit]

domain (plural domains)

1. A geographic area owned or controlled by a single person or organization. [quotations ▾]
The king ruled his **domain** harshly.
2. A field or sphere of activity, influence or expertise.
Dealing with complaints isn't really my **domain**: get in touch with customer services.
His **domain** is English history.
3. A group of related items, topics, or subjects. [quotations ▾]
4. (mathematics) The set of all possible mathematical entities (points) where a given function is defined.
5. (mathematics, set theory) The set of input (argument) values for which a function is defined.
6. (mathematics) A ring with no zero divisors; that is, in which no product of nonzero elements is zero.
Hyponym: integral domain
7. (mathematics, topology, mathematical analysis) An open and connected set in some topology. For example, the interval (0,1) as a subset of the real numbers.
8. (computing, Internet) Any DNS domain name, particularly one which has been delegated and has become representative of the delegated domain name and its subdomains. [quotations ▾]
9. (computing, Internet) A collection of DNS or DNS-like domain names consisting of a delegated domain name and all its subdomains.
10. (computing) A collection of information having to do with a domain, the computers named in the domain, and the network on which the computers named in the domain reside.
11. (computing) The collection of computers identified by a domain's domain names.
12. (physics) A small region of a magnetic material with a consistent magnetization direction.
13. (computing) Such a region used as a data storage element in a bubble memory.
14. (data processing) A form of technical metadata that represent the type of a data item, its characteristics, name, and usage. [quotations ▾]
15. (taxonomy) The highest rank in the classification of organisms, above kingdom; in the three-domain system, one of the taxa *Bacteria*, *Archaea*, or *Eukaryota*.
16. (biochemistry) A folded section of a protein molecule that has a discrete function; the equivalent section of a chromosome

Figure 1- 7. Domain: Means many things depending on context

In the context of a business however, the word domain covers the overall scope of its primary activity—the service it provides to its customers. This is also referred as the **problem domain**. For example, Tesla operates in the domain of electric vehicles, Netflix provides online movies and shows, while McDonald's provides fast food. Some companies like Amazon, provide services in more than one domain—online retail, cloud computing, among others.

1.4.2. What is a Subdomain?

The domain of a business (at least the successful ones) almost always encompasses fairly complex and abstract concepts. With a view to better deal with this complexity, domain-driven design advises decomposing the domain of a business into multiple manageable parts called **subdomains**. This facilitates better understanding and makes it easier to arrive at a solution. For example, the online retail domain may be divided into subdomains such as product, inventory, rewards, shopping cart, order management, payments, shipping, etc. as shown below:

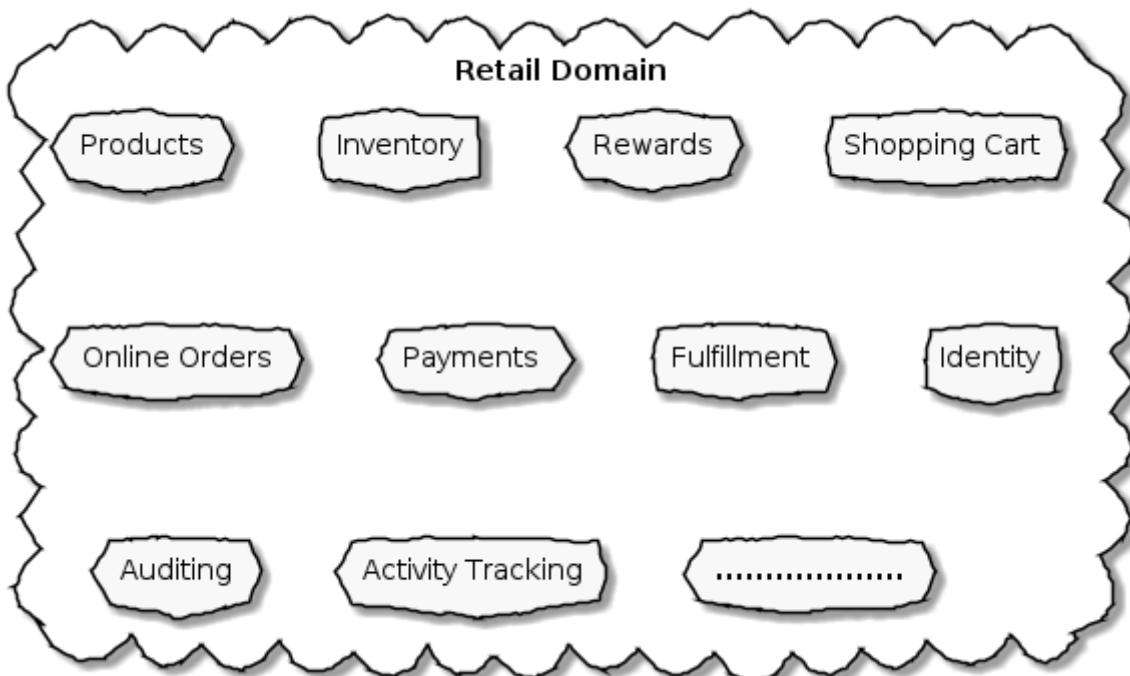


Figure 1- 8. Subdomains in the Retail domain

In certain businesses, subdomains themselves may turn out to become very complex on their own

and may require further decomposition. For instance, in the retail example above, it may be required to break the products subdomain into further constituent subdomains such as catalog, search, recommendations, reviews, etc. as shown below:

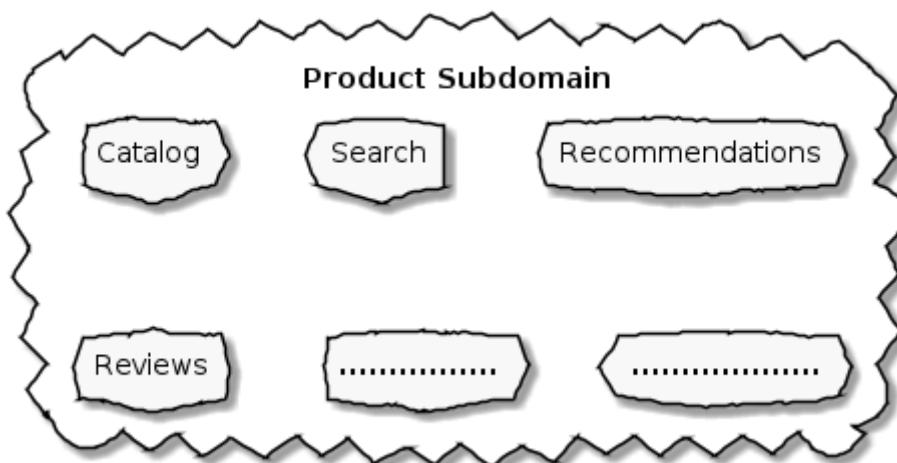


Figure 1- 9. Subdomains in the Products subdomain

Further breakdown of subdomains may be needed until we reach a level of manageable complexity.

1.4.3. Types of Subdomains

Breaking down a complex domain into more manageable subdomains is a great thing to do. However, not all subdomains are created equal. With any business, the following three types of subdomains are going to be encountered:

- **Core:** The main focus area for the business. This is what provides the biggest differentiation and value. It is therefore natural to want to place the most focus on the core subdomain. In the retail example above, shopping cart and orders might be the biggest differentiation — and hence may form the core subdomains for that business venture. It is prudent to implement core subdomains in-house given that it is something that businesses will desire to have the most control over. In the online retail example above,
- **Supporting:** Like with every great movie, where it is not possible to create a masterpiece without a solid supporting cast, so it is with supporting or auxiliary subdomains. Supporting subdomains are usually very important and very much required, but may not be the primary focus to run the business. These supporting subdomains, while necessary to run the business, do not usually offer a significant competitive advantage. Hence it might be even fine to completely outsource this work or use an off-the-shelf solution as is or with minor tweaks. For the retail example above, assuming that online ordering is the primary focus of this business, catalog management may be a supporting subdomain.
- **Generic:** When working with business applications, one is required to provide a set of capabilities **not** directly related to the problem being solved. Consequently, it might suffice to just make use of an off-the-shelf solution. For the retail example above, the identity, auditing and activity tracking subdomains might fall in that category.



It is important to note that the notion of core vs. supporting vs. generic subdomains is very context specific. What is core for one business may be supporting or generic for another. Identifying and distilling the core domain requires deep understanding and experience of what problem is being attempted to be solved.

1.4.4. Domain Experts

To run a successful digital business, you need specialists — those who have a deep and intimate understanding of the domain. Domain experts are subject matter experts (SMEs) who have a very strong grasp of the business. Domain experts may have varying degrees of expertise. Some SMEs may choose to specialize in specific subdomains, while others may have a broader understanding of how the overall business works.

Any modern software team requires expertise in at least two areas — the functionality of the domain and the art of translating it into high quality software. While the domain experts specify the **why** and the **what**, technical experts (software developers) specify the **how**. Strong contributions and synergy between both groups is absolutely essential to ensure sustained high performance and success.

1.4.5. Promoting a Shared Understanding

Previously, we saw how [organizational silos](#) can result in valuable information getting diluted. At a credit card company I used to work with, the words plastic, payment instrument, account, PAN (Primary Account Number), BIN (Bank Identification Number), card were all used by different team members to mean the exact same thing - the **credit card** when working in the same area of the application. To make matters worse, a lot of this muddled use of terms got implemented in code as well. While this might feel like a trivial thing, it had far-reaching consequences. Product experts, architects, developers, all came and went, each regressively contributing to more confusion, muddled designs, implementation and technical debt with every new enhancement — accelerating the journey towards the dreaded, unmaintainable, [big ball of mud](#).

DDD advocates breaking down these artificial barriers, and putting the domain experts and the developers on the same level footing by working collaboratively towards creating what DDD calls a **ubiquitous language** — a shared vocabulary of terms, words, phrases to continuously enhance the collective understanding of the entire team. This phraseology is then used actively in every aspect of the solution: the everyday vocabulary, the designs, the code — in short by **everyone** and **everywhere**. Consistent use of the common ubiquitous language helps reinforce a shared understanding and produce solutions that better reflect the mental model of the domain experts.

1.4.6. Evolving a Domain Model and a Solution

The ubiquitous language helps establish a consistent albeit informal lingo among team members. To enhance understanding, this can be further refined into a formal set of abstractions — a **domain model** to represent the solution in software. It is very important to note that this domain model is modeled to fall within the context of a single subdomain for which a solution is being explored, not the entire domain of the business. This solution is termed as a **bounded context** i.e. the ubiquitous language and domain model are only valid within those bounds and context — not outside of it.

This means that the system as a whole can be represented as a set of bounded contexts which have relationships with each other. These relationships define how these bounded contexts can integrate with each other and are called ***context maps***.

Care should be taken to retain focus on solving the business problem at hand at all times. Teams will be better served if they expend the same amount of effort modeling business logic as the technical aspects of the solution. To keep accidental complexity in check, it will be best to isolate the infrastructure aspects of the solution from this model. These models can take several forms, including conversations, whiteboard sessions, documentation, diagrams, tests and other forms of architecture fitness functions. It is also important to note that this is **not** a one-time activity. As the business evolves, the domain model and the solution will need to keep up. This can only be achieved through close collaboration between the domain experts and the developers at all times.

DDD has a catalog of strategic and tactical patterns which accelerate this process of continuous learning. In addition, modern techniques such as [domain storytelling](#), [event storming](#), and [evolutionary architecture](#) can greatly aid this process of evolving the ubiquitous language and domain model. We will examine all of these in much detail in upcoming chapters,

The thrust of DDD is that **one single model** form the bedrock of team communication, design, and implementation. While teams may and will indeed require a variety of means to express the model, it is very important to keep the executable code and the various representations up to date at all times.

1.4.7. The Essence of DDD

In this section we have taken a look at DDD at a very high level. Enclosed below is an attempt to capture the essence of what domain-driven design means.

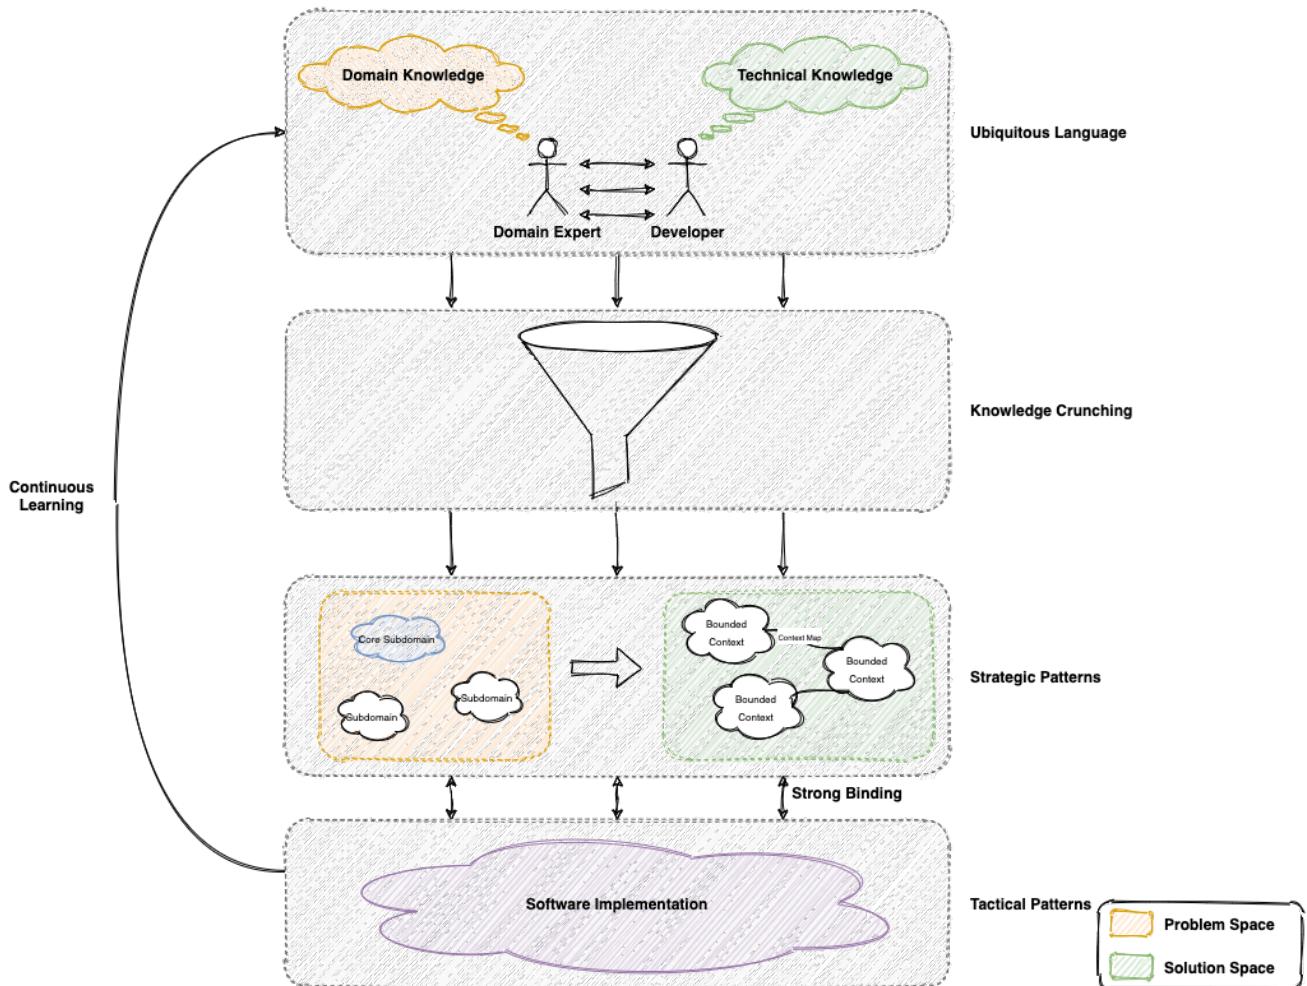


Figure 1- 10. Essence of DDD

In subsequent chapters we will reinforce all of the concepts introduced here in a lot more detail. In the next section, we will look at why the ideas of DDD, introduced all those years ago, are still very relevant. If anything, we will look at why they are becoming even more relevant now than ever.

1.5. Why is DDD Relevant? Why Now?

He who has a why to live for can bear almost any how.

— Friedrich Nietzsche

In a lot of ways, domain-driven design was way ahead of its time when Eric Evans introduced the concepts and principles back in 2003. DDD seems to have gone from strength to strength. In this section, we will examine why DDD is even more relevant today, than it was when Eric Evans wrote his book on the subject way back in 2003.

1.5.1. Rise of Open Source

Eric Evans, during his keynote address at the Explore DDD conference in 2017, lamented about how difficult it was to implement even the simplest concepts like immutability in value objects when his book had released. In contrast though, nowadays, it's simply a matter of importing a mature, well documented, tested library like [Project Lombok](#) or [Immutables](#) to be productive, literally in a matter of minutes. To say that open source software has revolutionized the software industry

would be an understatement! At the time of this writing, the public maven repository (<https://mvnrepository.com>) indexes no less than a staggering **18.3 million artifacts** in a large assortment of popular categories ranging from databases, language runtimes to test frameworks and many many more as shown in the chart below:

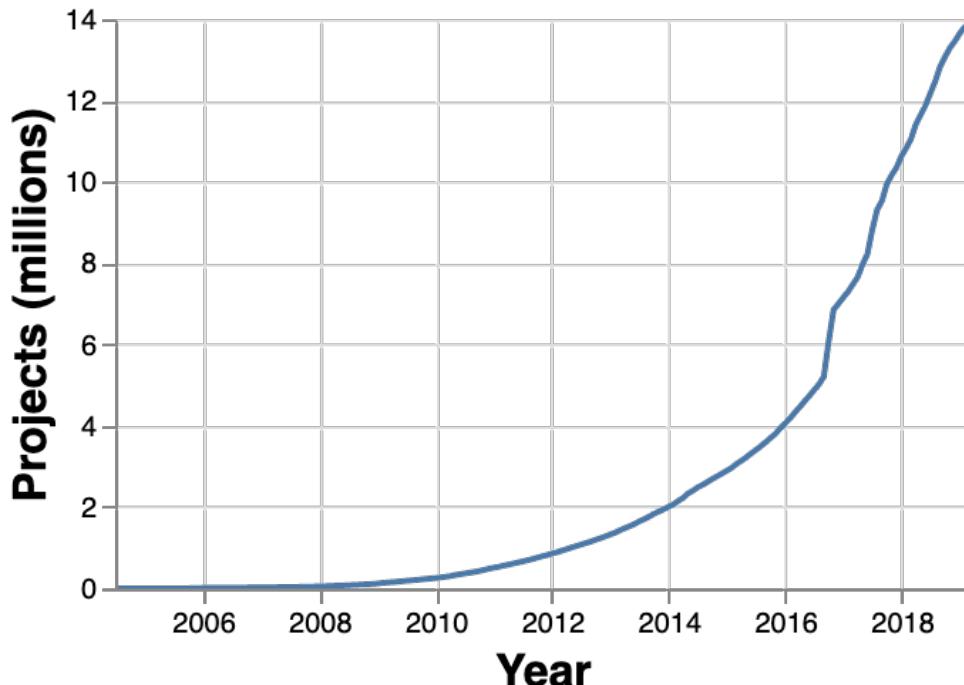


Figure 1- 11. Open source Java over the years. Source: <https://mvnrepository.com/>

Java stalwarts like the [spring framework](#) and more recent innovations like [spring boot](#), [quarkus](#), etc. make it a no-brainer to create production grade applications, literally in a matter of minutes. Furthermore, frameworks like [Axon](#), [Lagom](#), etc. make it relatively simple to implement advanced architecture patterns such as CQRS, event sourcing, that are very complementary to implementing DDD-based solutions.

1.5.2. Advances in Technology

DDD by no means is just about technology, it could not be completely agnostic to the choices available at the time. 2003 was the heyday of heavyweight and ceremony-heavy frameworks like J2EE (Java 2 Enterprise Edition), EJBs (Enterprise JavaBeans), SQL databases, ORMs (Object Relational Mappers) and the like—with not much choice beyond that when it came to enterprise tools and patterns to build complex software—at least out in the public domain. The software world has evolved and come a very long way from there. In fact, modern game changers like Ruby on Rails and the public cloud were just getting released. In contrast though, we now have no shortage of application frameworks, NoSQL databases, programmatic APIs to create infrastructure components with a lot more releasing with monotonous regularity.

All these innovations allow for rapid experimentation, continuous learning and iteration at pace. These game changing advances in technology have also coincided with the exponential rise of the internet and ecommerce as viable means to carry out successful businesses. In fact the impact of the internet is so pervasive that it is almost inconceivable to launch businesses without a digital component being an integral component. Finally, the consumerization and wide scale penetration of smartphones, IoT devices and social media has meant that data is being produced at rates

inconceivable as recent as a decade ago. This means that we are building for and solving the most complicated problems by several orders of magnitude.

1.5.3. Rise of Distributed Computing

There was a time when building large monoliths was very much the default. But an exponential rise in computing technology, public cloud, (IaaS, PaaS, SaaS, FaaS), big data storage and processing volumes, which has coincided with an arguable slowdown in the ability to continue creating faster CPUs, have all meant a turn towards more decentralized methods of solving problems.

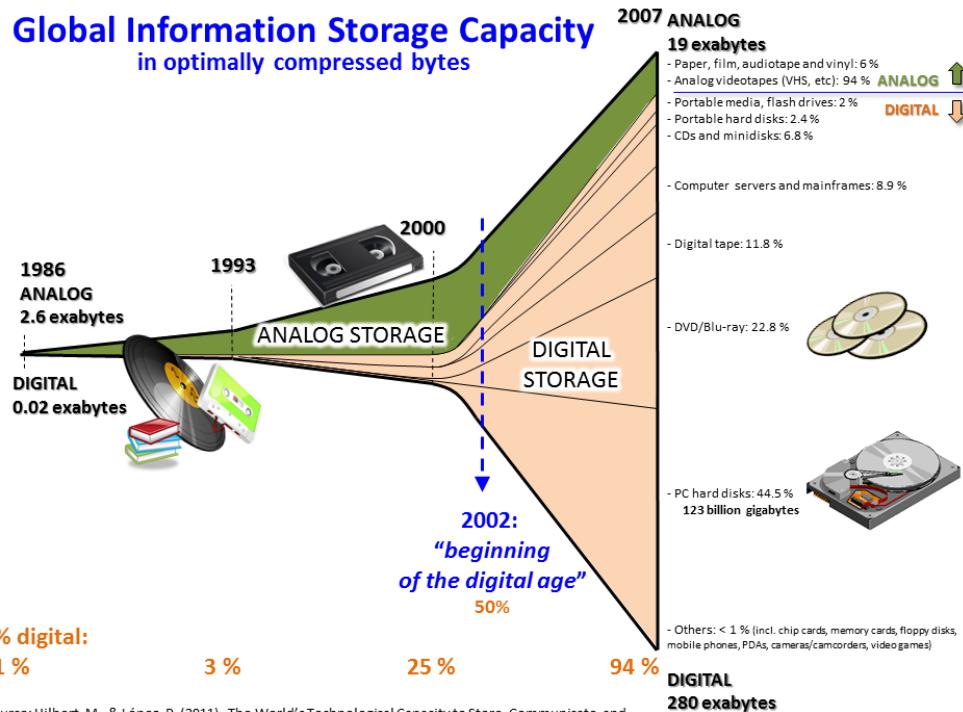


Figure 1- 12. Global Information Storage Capacity

Domain-driven design with its emphasis on dealing with complexity by breaking unwieldy monoliths into more manageable units in the form of subdomains and bounded contexts, fits naturally to this style of programming. Hence it is no surprise to see a renewed interest in adopting DDD principles and techniques when crafting modern solutions. To quote Eric Evans, it is no surprise that Domain-Driven Design is even more relevant now than when it was originally conceived!

1.6. Summary

In this chapter we examined some common reasons for why software projects fail. We saw how inaccurate or misinterpreted requirements, architecture (or the lack thereof), excessive technical debt, etc. can get in the way of meeting business goals and success.

We looked at the basic building blocks of domain-driven design such as domains, subdomains, ubiquitous language, domain models, bounded contexts and context maps. We also examined why the principles and techniques of domain-driven design are still very much relevant in the modern age of microservices and serverless. You should now be able to appreciate the basic terms of DDD and understand why it is important in today's context.

In the next chapter we will take a closer look at the real-world mechanics of domain-driven design. We will delve deeper into the strategic and tactical design elements of DDD and look at how using these can help form the basis for better communication and create more robust designs.

1.7. Questions

1. What are the most common reasons for software projects to fail?
2. What do the terms domain and sub-domain mean?
3. What are the different types of sub-domains?
4. What is the difference between sub-domains and bounded contexts?
5. Why is DDD relevant in today's context?

1.8. Further Reading

Title	Author	Location
Pulse of the Profession - 2017	PMI	https://www.pmi.org/-/media/pmi/documents/public/pdf/learning/thought-leadership/pulse/pulse-of-the-profession-2017.pdf
Pulse of the Profession - 2020	PMI	https://www.pmi.org/learning/library/forging-future-focused-culture-11908
Project success: Definitions and Measurement Techniques	PMI	https://www.pmi.org/learning/library/project-success-definitions-measurement-techniques-5460
Project success: definitions and measurement techniques	JK Pinto, DP Slevin	https://www.pmi.org/learning/library/project-success-definitions-measurement-techniques-5460
Analysis Paralysis	Ward Cunningham	https://proxy.c2.com/cgi/wiki?AnalysisParalysis
Big Design Upfront	Ward Cunningham	https://wiki.c2.com/?BigDesignUpFront
Enterprise Modeling Anti-Patterns	Scott W. Ambler	http://agilemodeling.com/essays/enterpriseModelingAntiPatterns.htm
A Project Manager's Guide To 42 Agile Methodologies	Henny Portman	https://thedigitalprojectmanager.com/agile-methodologies
Domain-Driven Design Even More Relevant Now	Eric Evans	https://www.youtube.com/watch?v=kIKwPNKXaLU
Introducing Deliberate Discovery	Dan North	https://dannorth.net/2010/08/30/introducing-deliberate-discovery/

Title	Author	Location
No Silver Bullet—Essence and Accident in Software Engineering	Fred Brooks	http://faculty.salisbury.edu/~xswang/Research/Papers/SERelated/no-silver-bullet.pdf
Mastering Non-Functional Requirements	Sameer Paradkar	https://www.packtpub.com/product/mastering-non-functional-requirements/9781788299237
Big Ball Of Mud	Brian Foote & Joseph Yoder	http://www.laputan.org/mud/
The Forgotten Layer of the Test Automation Pyramid	Mike Cohn	https://www.mountaingoatsoftware.com/blog/the-forgotten-layer-of-the-test-automation-pyramid
Tech debt: Reclaiming tech equity	Vishal Dalal et al	https://www.mckinsey.com/business-functions/mckinsey-digital/our-insights/tech-debt-reclaiming-tech-equity
Is High Quality Software Worth the Cost	Martin Fowler	https://martinfowler.com/articles/is-quality-worth-cost.html#WeAreUsedToATrade-offBetweenQualityAndCost

1.9. Answers

1. Refer to section 1.2
2. Refer to sections 1.4.1 and 1.4.2
3. Refer to section 1.4.3
4. Refer to section 1.4.7
5. Refer to section 1.5

Chapter 2. The Mechanics of Domain-Driven Design (30 Pages)

When eating an elephant, take one bite at a time.

— Creighton Abrams

As mentioned in the previous chapter, many things can put a project off course. In this chapter, we look at how DDD gives us a set of tenets and techniques to arrive at a collective understanding of the problem at hand in the face of ambiguity, break it down into manageable chunks and translate it into reliably working software.

2.1. Understanding the Problem

2.1.1. Problem Space vs. Solution Space

2.1.2. Dealing with Ambiguity

2.2. Arriving at a Shared Understanding

2.3. Breaking Down the Problem

2.3.1. What is a Domain?

2.3.2. What is a Sub-Domain?

2.3.3. The Core Sub-Domain

2.4. Modeling a Solution

2.4.1. What is a Model?



Anemic domain models

2.4.2. Context Maps

2.4.3. Bounded Contexts

2.5. Implementing the Solution

Chapter 3. Where and How Does DDD Fit? (15 pages)

We won't be distracted by comparison if we are captivated with purpose.

— Bob Goff

Software architecture refers to the fundamental structures of a software system and the discipline of creating such structures and systems. Over the years, we have accumulated a series of architecture styles to help us deal with system complexity. In this chapter we will examine how DDD compares with several of these architecture styles and how/where it fits in the overall scheme of things when crafting a software solution.

3.1. Architecture Styles

3.2. Layered Architecture

3.3. Onion Architecture

3.4. Hexagonal Architecture

3.5. Service Oriented Architecture

3.6. Microservice Architecture

3.7. Event-Driven Architecture (EDA)

3.8. Command Query Responsibility Segregation (CQRS)

3.9. Serverless Architecture

Part 2: Implementing DDD in the Real World

In Part 2, we will implement a real-world application using JVM-based technologies such as Vaadin, Spring Boot, Axon Framework, Cadence, among others.

Chapter 4. Domain Analysis and Modeling

He who asks a question remains a fool for five minutes. He who does not ask remains a fool forever.

— Chinese Proverb

4.1. Introduction

As we saw in the previous chapter, misinterpreted requirements cause a significant portion of software projects to fail. Arriving at a shared understanding and creating a useful domain model necessitates high degrees of collaboration with domain experts. In this chapter, we will introduce the sample application we will use throughout the book and explore modeling techniques such as domain storytelling and eventstorming to enhance our collective understanding of the problem in a reliable and structured manner.

The following topics will be covered in this chapter:

- Introducing the example application (Letter of Credit)
- Enhancing shared understanding
- Domain storytelling
- EventStorming

This chapter will help developers and architects learn how to apply these techniques in real-life situations to produce elegant software solutions that mirror the problem domain that needs to be solved. Similarly, non-technical domain experts will understand how to communicate their ideas and collaborate effectively with technical team members to accelerate the process of arriving at a shared understanding.

4.2. Technical requirements

There are no specific technical requirements for this chapter. However, given that it may become necessary to collaborate remotely as opposed to being in the same room with access to a whiteboard, it will be useful to have access to the following:

1. Digital whiteboard (like <https://www.mural.co/> or <http://miro.com/>)
2. Online domain storytelling modeler (like <https://www.wps.de/modeler/>)

4.3. Introducing the LC application

In many countries, international trade represents a significant portion of the gross domestic product (GDP)—making an exchange of capital, goods, and services between untrusted parties spread across the globe a necessity. While economic organizations such as the World Trade Organization (WTO) were formed specifically to ease and facilitate this process, differences in factors such as economic policy, trade laws, currency, etc. ensure that carrying out trade

internationally can be a complex process with several entities involved across countries. Letter of Credit exist to simplify this process. Let's take a look at how they work.

4.3.1. What is a Letter of Credit (LC)

Documentary Letter of Credit (LC) is a financial instrument issued by the banks as a contract between the importer (or buyer) and the exporter (or seller). This contract specifies terms and conditions of the transaction under which importer promises to pay the exporter in exchange for the goods or services provided by the exporter. Letter of Credit transaction typically involves multiple parties. A simplified summary of the parties involved is described below:

1. **Importer:** The buyer of the goods or services.
2. **Exporter:** The seller of the goods or services.
3. **Freight Forwarder:** The agency that handles shipment of goods on behalf of the exporter. This is only applicable in cases there is an exchange of physical goods.
4. **Issuing Bank:** The bank that the importer requests to issue the LC application. Usually the importer has a pre-existing relationship with this bank.
5. **Advising Bank:** The bank that informs the exporter about the issuance of the LC. This is usually a bank that is native to the exporter's country.
6. **Negotiating Bank:** The bank that the exporter submits documents for the shipment of goods, or the services provided. Usually the exporter has a pre-existing relationship with this bank.
7. **Reimbursement Bank:** The bank that reimburses the funds to the negotiating bank, at the request of the issuing bank.



It is important to note that the same bank can play more than one role for a given transaction. In the most complex cases, there can be four distinct banks involved for a transaction (sometimes even more, but we will skip those cases for brevity).

4.3.2. The LC issuance application

XYZ Bank has reached out to us in order to automate the process of LC application and issuance. In this chapter, and indeed the rest of this book, we will strive to understand, evolve, design and build a software solution to automate this process.

We understand that unless one is an expert dealing with international trade, it is unlikely that one would have an intimate understanding of concepts like Letters of Credit (LCs). In the upcoming section, we will look at demystifying LCs and how to work with them.

4.4. Enhancing shared understanding

When working with a problem where domain concepts are unclear, there is a need to arrive at a common understanding among key team members (both those that have bright ideas—the business/product people, and those that translate those ideas into working software—the software developers). For this process to be effective, we tend to look for approaches that are:

- Quick, informal and effective

- Collaborative - Easy to learn and adopt for both non-technical and technical team members
- Pictorial - because a picture can be worth a thousand words
- Usable for both coarse grained and fine-grained scenarios

There are several means to arrive at this shared understanding. Some commonly used approaches are listed below:

- UML
- BPMN
- Use Cases
- User Story Mapping
- CRC Models
- Data Flow Diagrams

Above modeling techniques have tried to formalize knowledge and express them in form of a structure diagram or text to help in delivering the business requirements as a software product. However, this attempt has not narrowed but has widened the gap between the business and, the software systems.

We will use **domain storytelling** and **eventstorming** as our means to capture business knowledge from domain experts for consumption of Developers, Business Analysts etc.

4.5. Domain storytelling

You're never going to kill storytelling because it's built into the human plan.
We come with it.

— Margaret Atwood

4.5.1. Introducing Domain Storytelling

Scientific research has now proven that learning methods that employ audio-visual aids assist both the teacher and the learners in retaining and internalizing concepts very effectively. In addition, teaching what one has learnt to someone else helps reinforce ideas and also stimulates the formation of new ones. Domain storytelling is a collaborative modeling technique that combines a pictorial language, real-world examples, and a workshop format to serve as a very simple, quick and effective technique for sharing knowledge among team members. Domain Storytelling is a technique invented and popularized by Stefan Hofer and Henning Schwentner based on some related work done at the University of Hamburg called *cooperation pictures*.

A pictorial notation of the technique is illustrated in the diagram below:

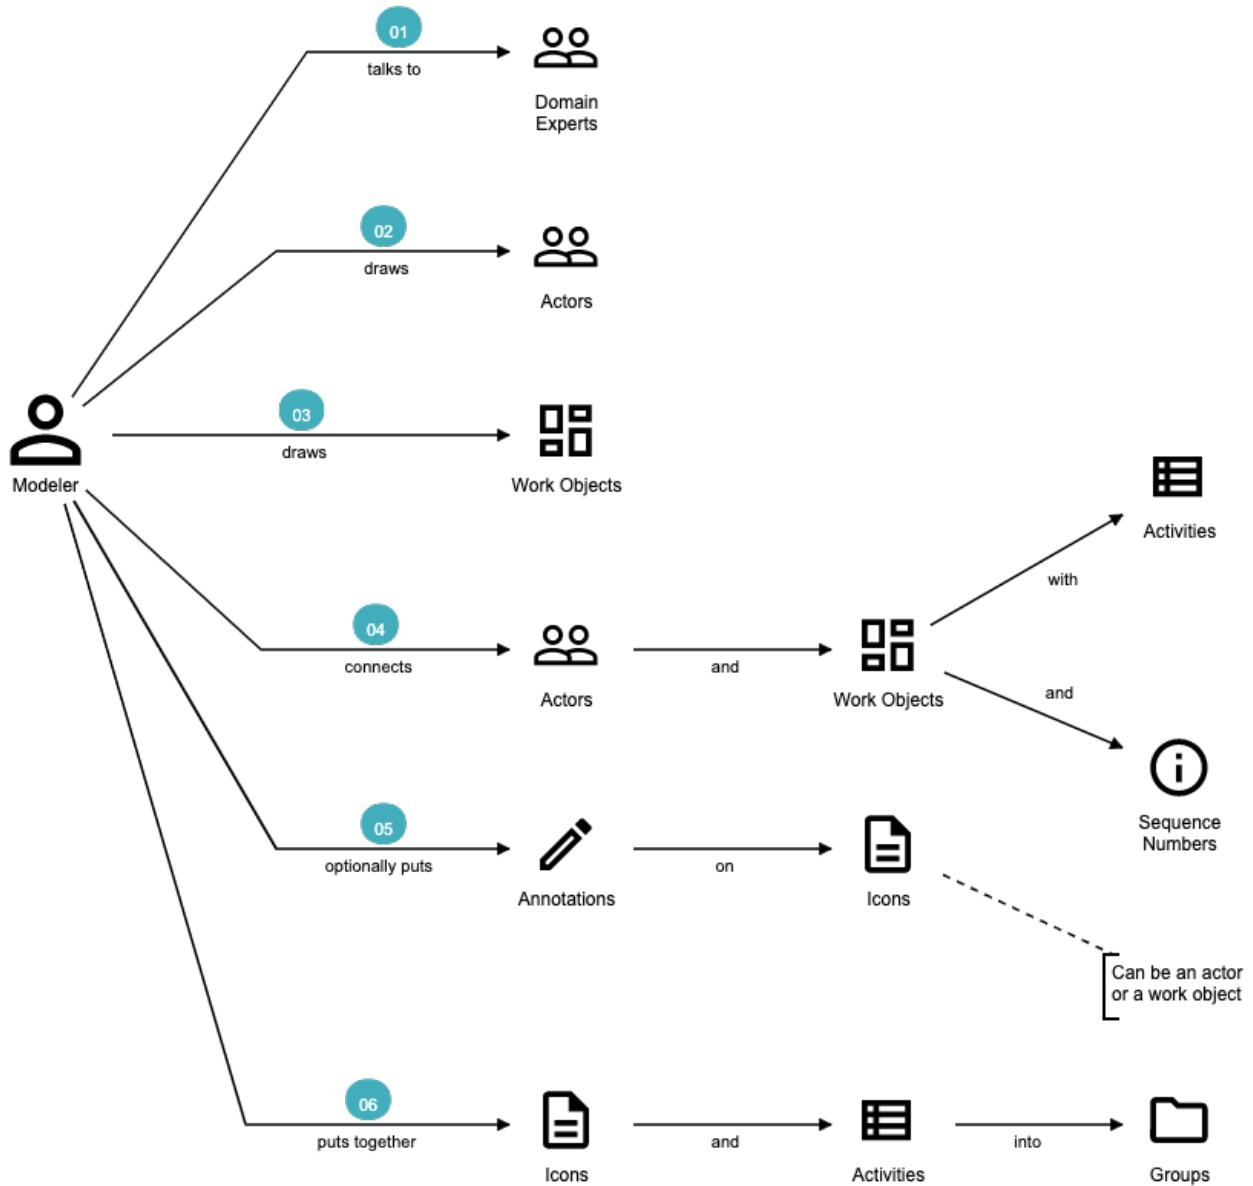


Figure 1- 13. Domain storytelling summarized

A domain story is conveyed using the following attributes:

Actors - Stories are communicated from the perspective of an actor (noun), for example, the issuing bank, who plays an active role in the context of that particular story. It is a good practice to use the ubiquitous language for the particular domain.

Work Objects - Actors act on some object, for example, applying for an LC. Again, this would be a term (noun) commonly used in the domain.

Activities - Actions (verb) performed by the actor on a work object. Represented by a labelled arrow connecting the actor and the work object.

Annotations - Used to capture additional information as part of the story, usually represented in few sentences.

Sequence Numbers - Usually, stories are told one sentence after the other. Sequence numbers helps capture the sequence of the activities in a story.

Groups - An outline to represent a collection of related concepts ranging from repeated/optional activities to sub-domains/organizational boundaries.

4.5.2. Using DST for the LC application

XYZ Bank has a process that allows processing of LCs. However, this process is very archaic, paper-based and manually intensive. Very few at the bank fully understand the process end-to-end and natural attrition has meant that the process is overly complex without good reason. So they are looking to digitize and simplify this process. DST itself is just a graphical notation which can be done in isolation. However, it is typical to not do this on your own and employ a workshop style with domain experts and software experts working collaboratively.

In this section, we will employ a DST workshop to capture the current business flow. The following is an excerpt of such a conversation between **Katie, the domain expert** and **Patrick, the software developer**.

Patrick : "Can you give me a high level overview of a typical LC Flow?"

Katie : "Sure, it all begins with the importer and the exporter entering into a contract for purchase of goods or services."

Patrick : "What form does this contract take? Is it a formal documentClause? Or is this just a conversation?"

Katie : "This is just a conversation."

Patrick : "Oh okay. What does the conversation cover?"

Katie : Several things — nature and quantity of goods, pricing details, payment terms, shipment costs and timelines, insurance, warranty, etc. These details may be captured in a purchase order—which is a simple documentClause elaborating the above.

At this time, Patrick draws this part of the interaction between the importer and the exporter. This graphic is depicted in the following diagram:



Figure 1- 14. Interaction between importer and exporter

Patrick : "Seems straight forward, so where does the bank come into the picture?"

Katie : "This is international trade and both the importer and the exporter need to mitigate the financial risk involved in such business transactions. So they involve a bank as a trusted mediator."

Patrick : "What kind of bank is this?"

Katie : "Usually, there are multiple banks involved. But it all starts with an **issuing bank**."

Patrick : "What is an issuing bank?"

Katie : "Any bank that is authorized to mediate international trade deals. This has to be a bank in the importer's country."

Patrick : "Does the importer need to have an existing relationship with this bank?"

Katie : "Not necessarily. There may be other banks with whom the importer may have a relationship with—which in turn liaises with the issuing bank on the importer's behalf. But to keep it simple, let's assume that the importer has an existing relationship with the issuing bank—which is our bank in this case."

Patrick : "Does the importer provide details of the purchase order to the issuing bank to get started?"

Katie : "Yes. The importer provides the details of the transaction by making an **LC application**."

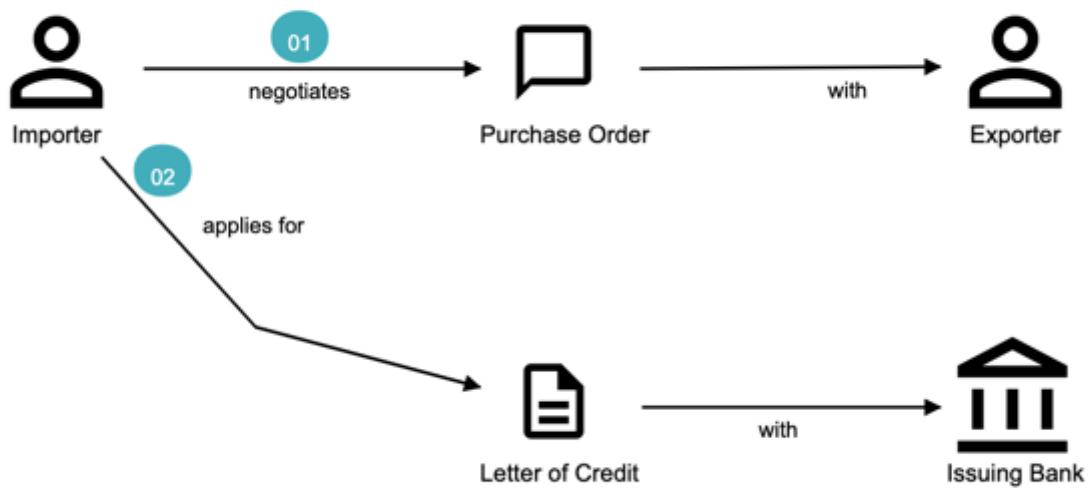


Figure 1- 15. Introducing the LC and the issuing bank

Patrick : "What does the issuing bank do when they receive this LC application?"

Katie : "Mainly two things—whether the financial standing of the importer and the legality of the goods being imported."

Patrick : "Okay. What happens if everything checks out?"

Katie : "The issuing bank approves the LC and notifies the importer."

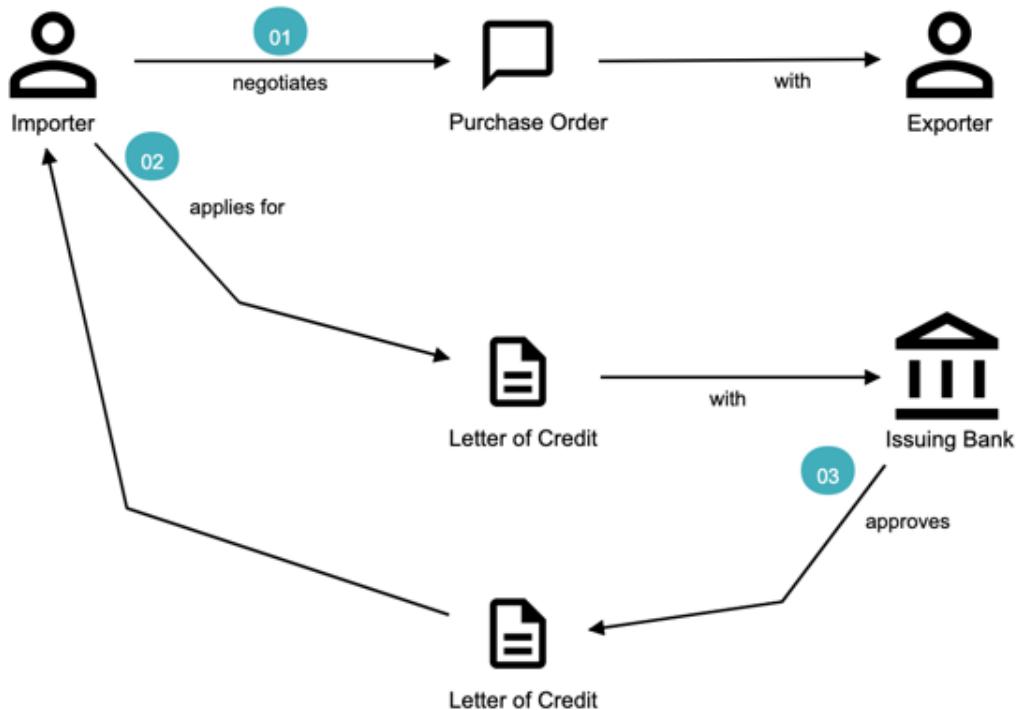


Figure 1- 16. Notifying LC approval to the importer

Patrick : "What happens next? Does the issuing bank contact the exporter now?"

Katie : "Not yet. It is not that simple. The issuing bank can only deal with a counterpart bank in the exporter's country. This bank is called the **advising bank**."

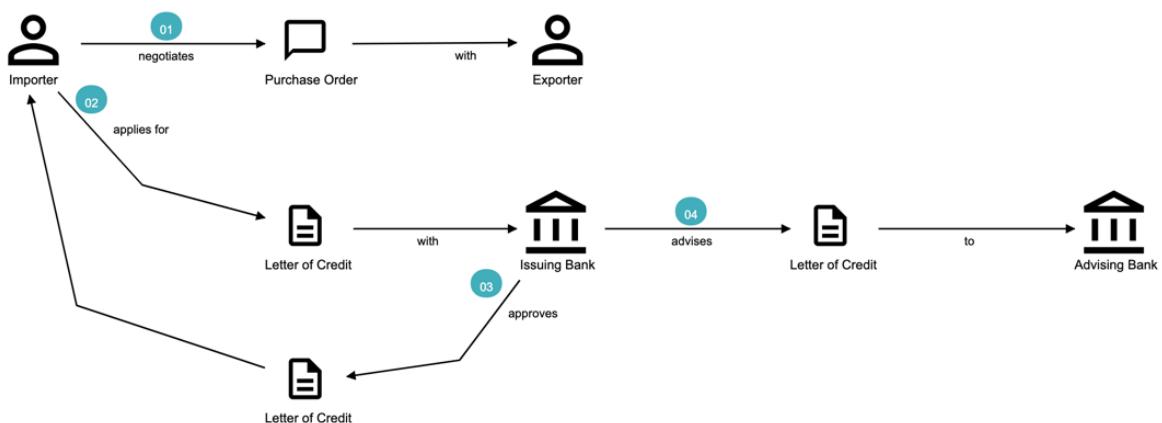


Figure 1- 17. Introducing the advising bank

Patrick : "What does the advising bank do?"

Katie : "The advising bank notifies the exporter about the LC."

Patrick : "Doesn't the importer need to know that the LC has been advised?"

Katie : "Yes. The issuing bank notifies the importer that the LC has been advised to the exporter."

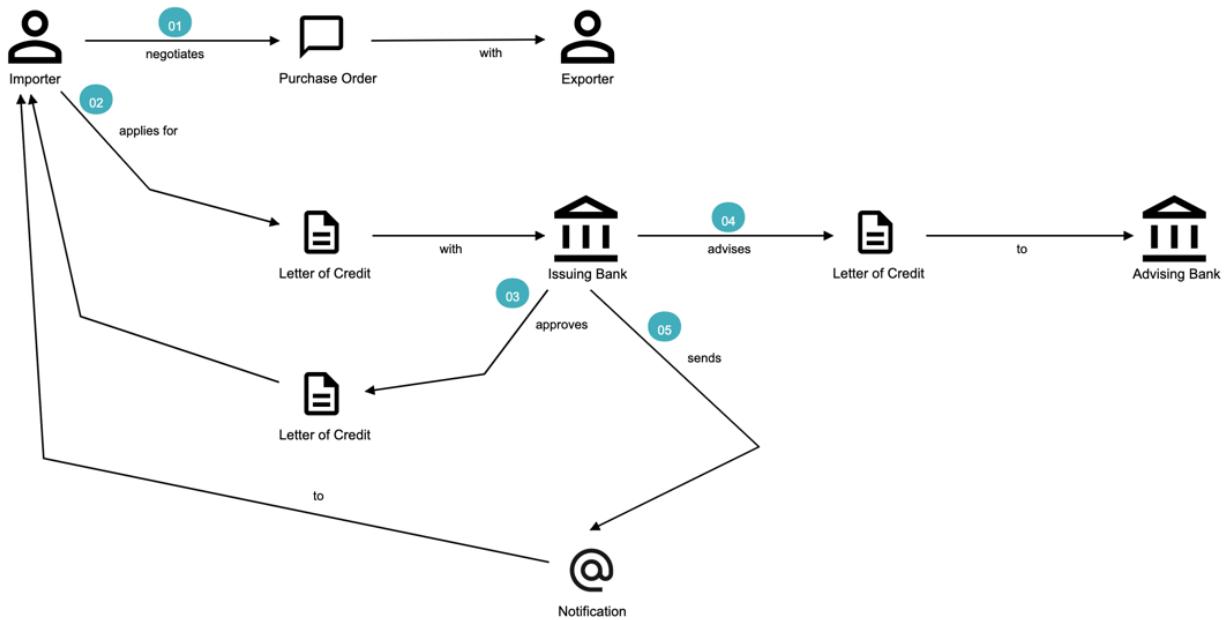


Figure 1- 18. Advice notification to the importer

Patrick : "How does the exporter know how to proceed?"

Katie : "Through the advising bank — they notify the exporter that the LC was issued."

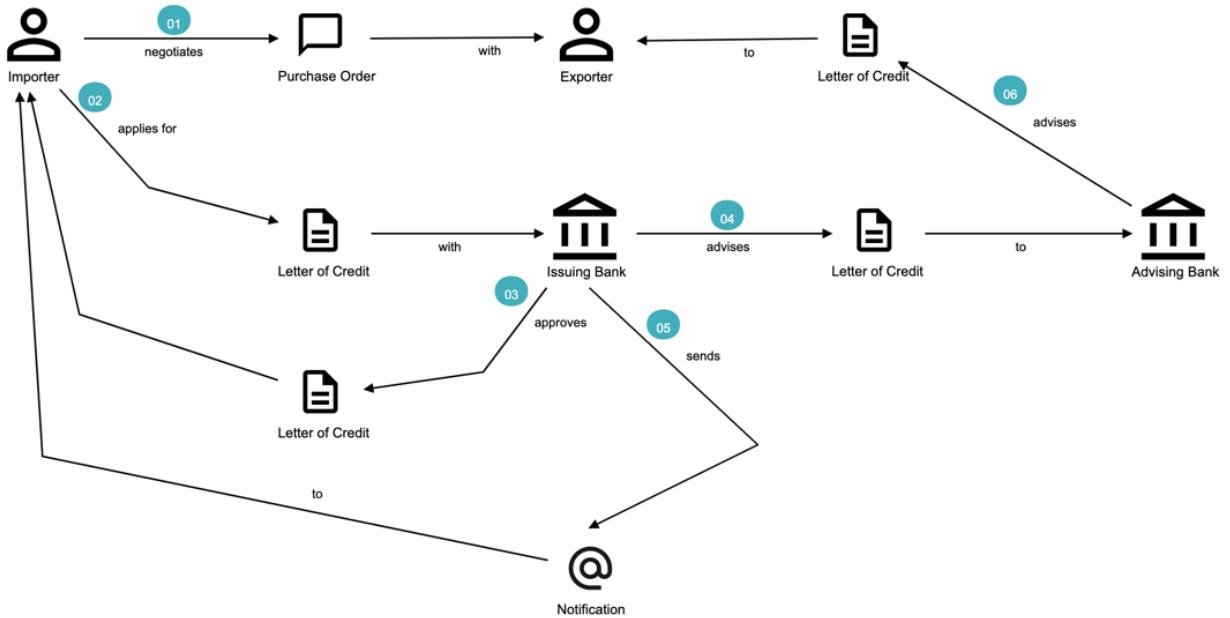


Figure 1- 19. Dispatching the advice to the exporter

Patrick : "Does the exporter initiate shipping at this time and how do they get paid?"

Katie : "Through the advising bank — they notify the exporter that the LC was issued and this triggers the next steps in the process — this process of settling the payment is called **settlement**. But let's focus on issuance right now. We will discuss settlement at a later time."

We have now looked at an excerpt of a typical DST workshop. The DST workshop has served to provide a reasonably good understanding of the high level business flow. Note that we have not referenced any technical artifacts during the process.

To be able to refine this flow and convert it into a form that can be used to design the software solution, we will need to further enhance this view. In the upcoming section, we will use EventStorming as a structured approach to achieve that.

4.6. EventStorming

The amount of energy necessary to refute bullshit is an order of magnitude bigger than to produce it.

— Alberto Brandolini

4.6.1. Introducing EventStorming

In the previous section, we gained a high level understanding of the LC Issuance process. To be able to build a real-world application, it will help to use a method that delves into the next level of detail. EventStorming, originally conceived by Alberto Brandolini, is one such method for the collaborative exploration of complex domains.

In this method, one simply starts by listing out all the events that are significant to the business domain in roughly chronological order on a wall or whiteboard using a bunch of colored sticky notes. Each of the note types (denoted by a color) serve a specific purpose as outlined below:

- **Domain Event:** An event that is significant to the business process — expressed in past tense.
- **Command:** An action or an activity that may result in one or more domain events occurring. This is either user initiated or system initiated, in response to a domain event.
- **User:** A person who performs a business action/activity.
- **Policy:** A set of business rules that need to be adhered to, for an action/activity to be successfully performed.
- **Read Model:** A piece of information required to perform an action/activity.
- **External System:** A system significant to the business process, but out of scope in the current context.
- **Hotspot:** Point of contention within the system that is likely confusing and/or puzzling beyond a small subsection of the team.
- **Aggregate:** An object graph whose state changes consistently and atomically.

The depiction of the stickies for our EventStorming workshop is shown here:

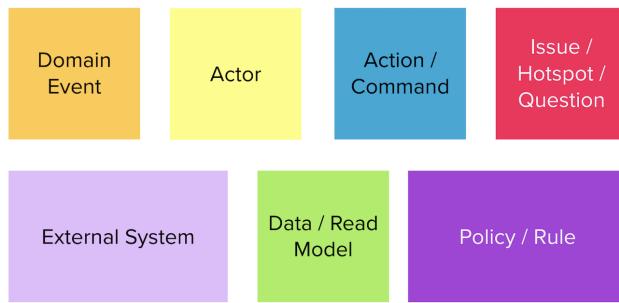


Figure 1- 20. EventStorming legend



Why domain events? When trying to understand a business process, it is convenient to express significant facts or things that happen in that context. It can also be informal and easy for audiences that are uninitiated with this practice. This provides an easy to digest visual representation of the domain complexity.

4.6.2. Using eventStorming for the LC issuance application

Now that we have a high level understanding of the current business process, thanks to the domain storytelling workshop, let's look at how we can delve deeper using eventstorming. The following is an excerpt of the stages from an eventstorming workshop for the same application.

1. Outline the event chronology

During this exercise, we recall significant **domain events** (using orange stickies) that happen in the system and paste them on the whiteboard, as depicted below. We ensure that the event stickies are pasted roughly in the chronological order of occurrence. As the timeline is enforced, the business flow will begin to emerge.



Figure 1- 21. Event chronology

This acts as an aid to understand the big picture. This also enables people in the room to identify hotspots in the existing business process. In the above illustration, we realized that, the process to handle "declined LC applications" is sub-optimal, i.e. applicants do not receive any information when their application is declined.

To address this, we added a new domain event which explicitly indicates that an application is

declined, as depicted below:



Figure 1- 22. New event to handle declined applications

2. Identify triggering activities and external systems

Having arrived at a high level understanding of event chronology, the next step is to embellish the visual with **activities/actions** that cause these events to occur (using blue stickies) and interactions with **external systems** (using pink stickies).

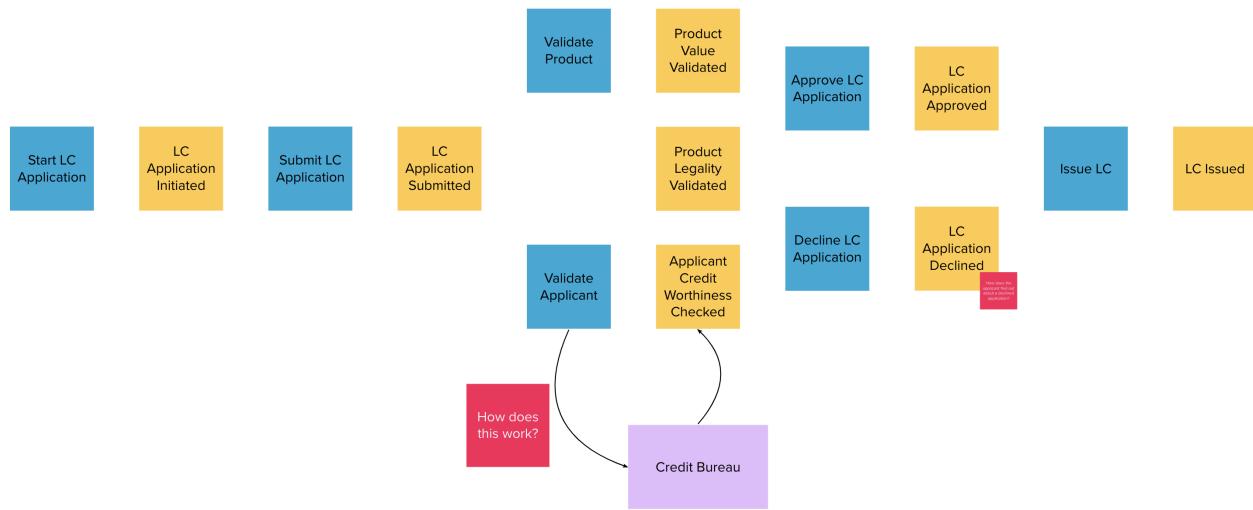


Figure 1- 23. Activities and external systems

3. Capture users, context and policies

The next step is to capture **users** who perform these activities along with their functional **context** (using yellow stickies) and policies (using purple stickies).

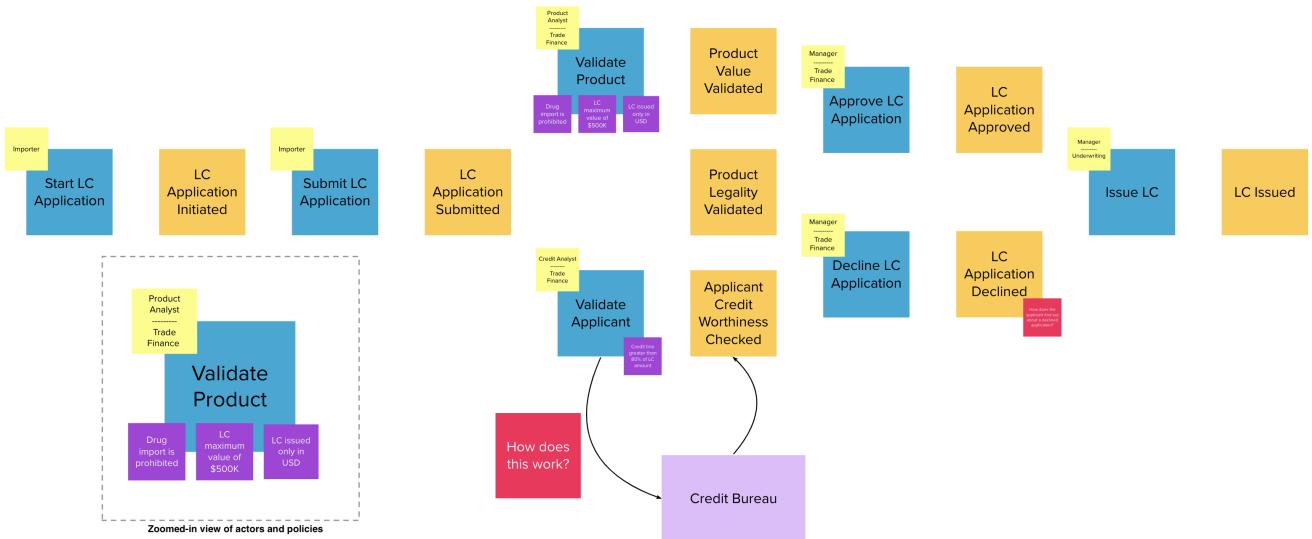


Figure 1-24. Users and policies

4. Outline read models

Every activity requires a certain set of data to be able to be performed. Users will need to view out-of-band data to that they need to act upon and also see the result of their actions. These sets of data are represented as **read models** (using green stickies).

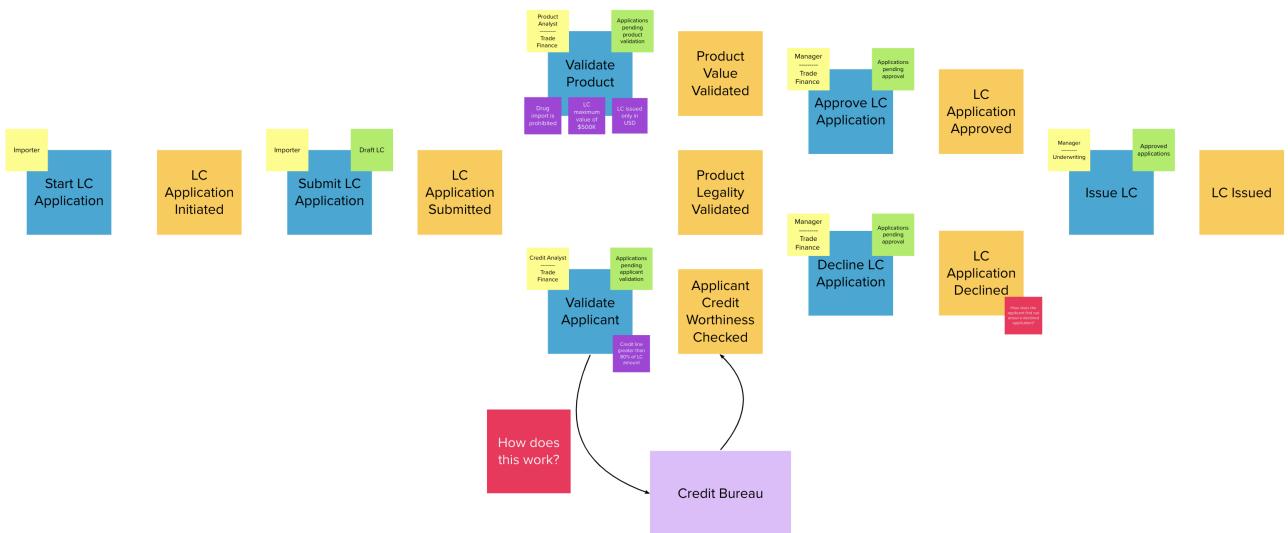


Figure 1-25. Big picture eventstorming workshop board



For both the domain storytelling and eventstorming workshops, it works best when we have approximately 6-8 people participating with a right mix of domain and technology experts.

This concludes the eventstorming workshop to gain a reasonably detailed understanding of the LC application and issuance process. Does this mean that we have concluded the domain requirements gathering process? Not at all—while we have made significant strides in understanding the domain, there is still a long way to go. The process of elaborating domain requirements is perpetual. Where are we in this continuum? The picture below is an attempt to clarify:

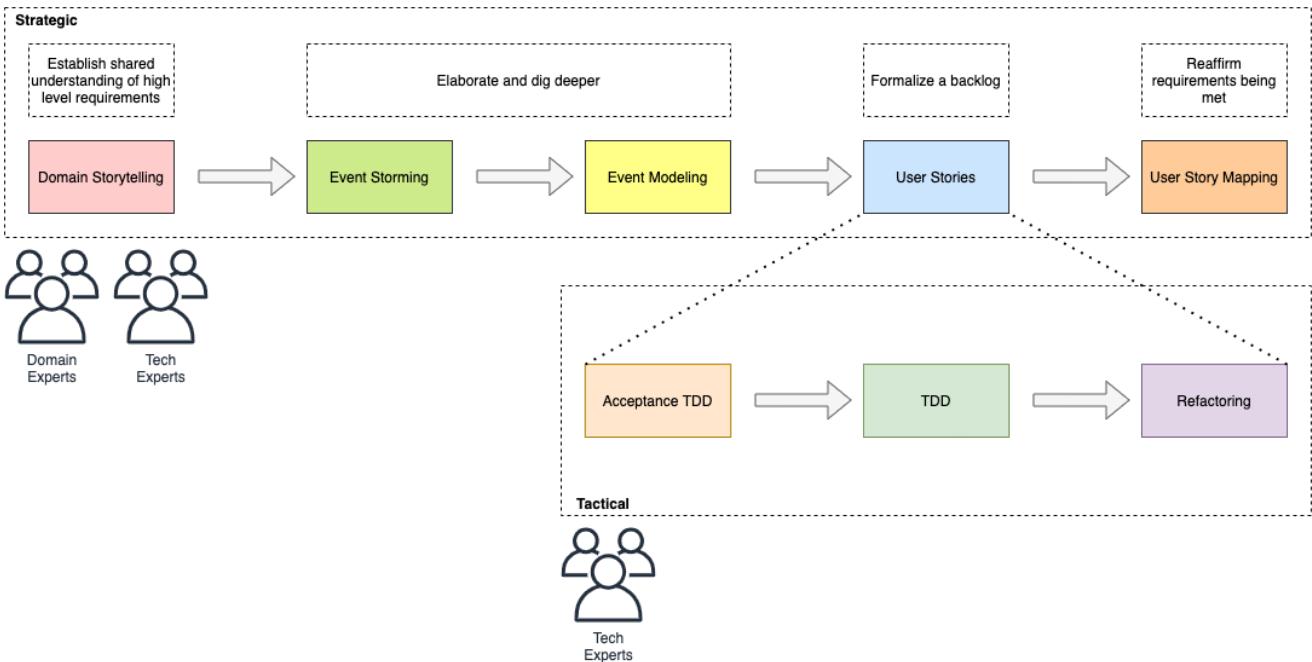


Figure 1- 26. Domain requirements elaboration continuum

In subsequent chapters we will examine the other techniques in more detail.

4.7. Summary

In this chapter we examined two ways to enhance our collective understanding of the problem domain using two lightweight modeling techniques — domain storytelling and eventstorming.

Domain storytelling uses a simple pictorial notation to share business knowledge among domain experts and technical team members. Eventstorming, on the other hand, uses a chronological ordering of domain events that occur as part of the business process to gain that same shared understanding.

Domain storytelling can be used as an introductory technique to establish high level understanding of the problem space, while eventstorming can be used to inform detailed design decisions of the solution space.

With this knowledge, we should be able to dive deeper into the technical aspects of solution implementation. In the next chapter, we will start implementation of the business logic, model our aggregate along with commands and domain events.

4.8. Questions

1. When should you use domain storytelling?
2. Pick an application in your current context. Can you use domain storytelling to capture actors, work objects and activities for the scenario you picked?
3. When should you use eventstorming?
4. Pick an application in your current context. Can you use eventstorming to capture domain events, actors, actions, hotspots, read models, external systems, etc. for the scenario you picked?

4.9. Further Reading

Title	Author	Location
Domain Storytelling	Stefan Hofer and Henning Schwentner	https://leanpub.com/domainstorytelling
An Introduction to Domain Storytelling	Virtual Domain-Driven Design	https://www.youtube.com/watch?v=d9k9Szkdprk
Domain Storytelling Resources	Stefan Hofer	https://github.com/hofstef/awesome-domain-storytelling
Introducing EventStorming	Alberto Brandolini	https://leanpub.com/introducing_eventstorming
Introducing Event Storming	Alberto Brandolini	https://ziobrando.blogspot.com/2013/11/introducing-event-storming.html
Event storming for fun and profit	Dan Terhorst-North	https://speakerdeck.com/tastapod/event-storming-for-fun-and-profit
EventStorming	Allen Holub	https://holub.com/event-storming/

4.10. Answers

1. Refer to section [Section 4.5.1](#)
2. Share and validate the domain storytelling artifact you created with your teammates.
3. Refer to section [Section 4.6.1](#)
4. Share and validate the eventstorming artifact you created with your teammates.

Chapter 5. Implementing Domain Logic

To communicate effectively, the code must be based on the same language used to write the requirements—the same language that the developers speak with each other and with domain experts.

— Eric Evans

In the Command Query Responsibility Segregation (CQRS) section, we describe how DDD and CQRS complement each other and how the command side (write requests) is the home of business logic. In this chapter, we will implement the command side API for the LC application using Spring Boot and Axon Framework, JSR-303 Bean Validations and persistence options by contrasting between state-stored vs event-sourced aggregates. The list of topics to be covered is as follows:

- Identifying aggregates
- Handling commands and emitting events
- Test-driving the application
- Persisting aggregates
- Performing validations

By the end of this chapter, you would have learnt how to implement the core of your system (the domain logic) in a robust, well encapsulated manner. You will also learn how to decouple your domain model from persistence concerns. Finally, you will be able to appreciate how to perform DDD's tactical design using services, repositories, aggregates, entities and value objects.

5.1. Technical Requirements

- JDK 1.8+ (We have used Java 16 to compile sample sources)
- Maven 3.x
- Spring Boot 2.4.x
- JUnit 5.7.x (Included with spring boot)
- Axon Framework 4.4.7 (DDD and CQRS Framework)
- Project Lombok (To reduce verbosity)
- Moneta 1.4.x (Money and currency reference implementation - JSR 354)

5.2. Command Query Responsibility Segregation (CQRS)

In this chapter, we will make use of the Command Query Responsibility Segregation (CQRS) architecture pattern to express the domain logic for our solution. The CQRS pattern strictly separates **write** (those that mutate state) and **read** (those that answer questions) operations.

5.2.1. Recap: What is CQRS

In traditional applications, a single domain, data/persistence model is used to handle all kinds of operations. With CQRS, we create distinct models to handle updates and enquiries. This is depicted in the following diagram:

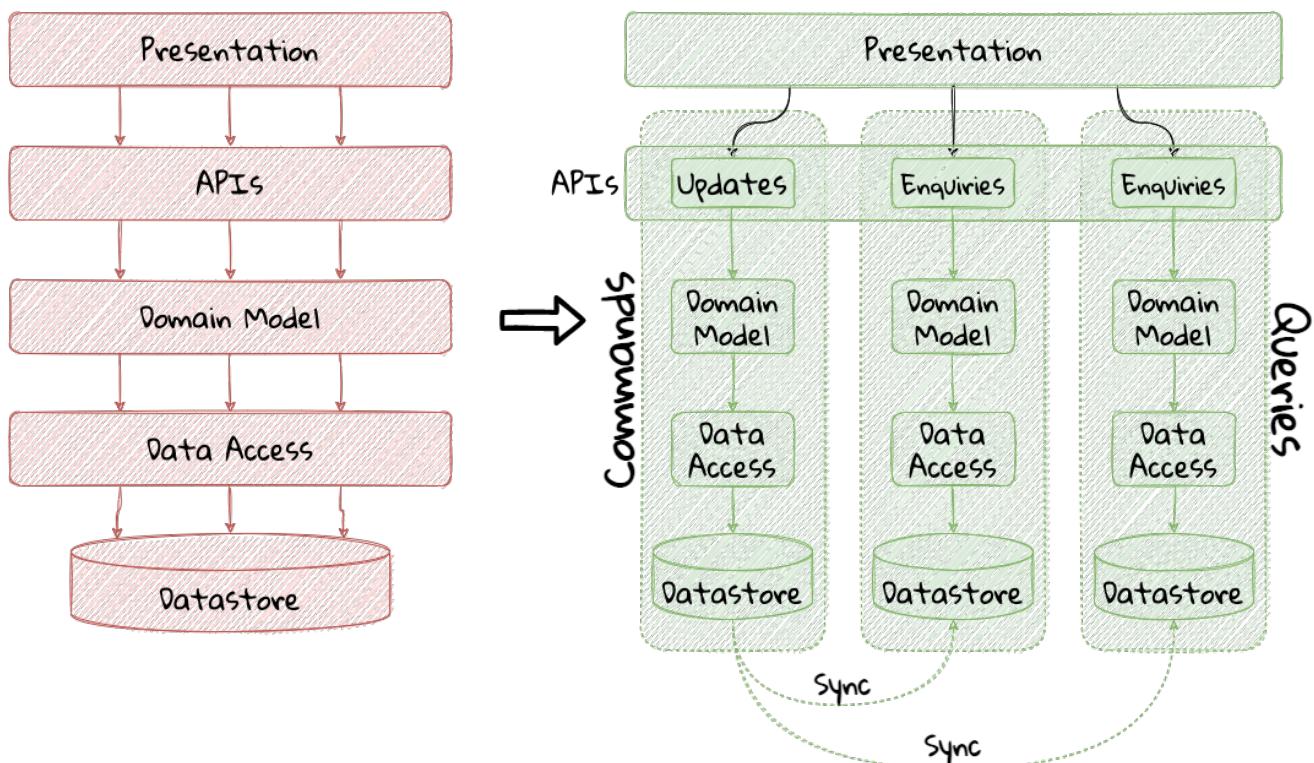


Figure 1-27. Traditional vs. CQRS Architecture



We depict multiple read models above because it is possible (but not necessary) to create more than one read model, depending on the kinds of query use cases that need to be supported.

For this to work predictably, the read model(s) need to be kept in sync with the write models (we will examine some of the techniques to do that in detail later).

5.2.2. Why CQRS?

The traditional, single-model approach works well for simple, CRUD-style applications, but starts to become unwieldy for more complex scenarios. We discuss some of these scenarios below:

- **Volume imbalance between read and writes:** In most systems, read operations often outnumber write operations by significant orders of magnitude. For example, consider the number of times a trader checks stock prices vs. the number of times they actually transact (buy or sell stock trades). It is also usually true that write operations are the ones that make businesses money. Having a single model for both reads and writes in a system with a majority of read operations can overwhelm a system to an extent where write performance can start getting affected.
- **Need for multiple read representations:** When working with relatively complex systems, it is not uncommon to require more than one representation of the same data. For example, when

looking at personal health data, one may want to look at a daily, weekly, monthly view. While these views can be computed on the fly from the *raw* data, each transformation (aggregation, summarization, etc.) adds to the cognitive load on the system. Several times, it is not possible to predict ahead of time, the nature of these requirements. By extension, it is not feasible to design a single canonical model that can provide answers to all these requirements. Creating domain models specifically designed to meet a focused set of requirements can be much easier.

- **Different security requirements:** Managing authorization and access requirements to data/APIs when working a single model can start to become cumbersome. For example, higher levels of security may be desirable for debit operations in comparison to balance enquiries. Having distinct models can considerably ease the complexity in designing fine-grained authorization controls.
- **More uniform distribution of complexity:** Having a single model to serve all use cases means that they can now be focused towards solving a single concern and thereby reduce complexity. It is worth noting that the essence of domain-driven design is mainly to work effectively with complex software systems and CQRS fits well with this line of thinking.

 Implementing CQRS does not require the use of any framework. However, in this book we will look at using Axon Framework because in our opinion it provides a set of conveniences to do so while staying out of the way. There are other frameworks that work comparably, like [Lagom Framework^{\[1\]}](#) and [Eventuate^{\[2\]}](#) that are worth exploring as well.

5.3. Bootstrapping the application

To get started, let's create a simple spring boot application using the following command:

```
curl -G https://start.spring.io/starter.zip \
      -d dependencies=web,data-jpa,lombok,validation,h2,actuator \
      -d name=lc-issuance-api \
      -d artifactId=lc-issuance-api \
      -d groupId=com.example.api \
      -d packaging=jar \
      -d description='LC Issuance API' \
      -d package-name=com.example.api \
      -o lc-issuance-api.zip
```

① The spring initializr to create the application archive in zip form

② The list of dependencies separated by a comma

③ The name of the archive containing the generated sources



Alternatively, you can use the spring initializr directly at <https://start.spring.io> or the spring boot CLI to generate the application.

This should create a file named `lc-issuance-api.zip` in the current directory. Unzip this file to a location of your choice and add a dependency on the Axon framework in the `dependencies` section of the `pom.xml` file:

```

<dependency>
    <groupId>org.axonframework</groupId>
    <artifactId>axon-spring-boot-starter</artifactId>
    <version>${axon-framework.version}</version> ①
</dependency>

```

① You may need to change the version

Also, add the following dependency on the **axon-test** library to enable unit testing of aggregates:

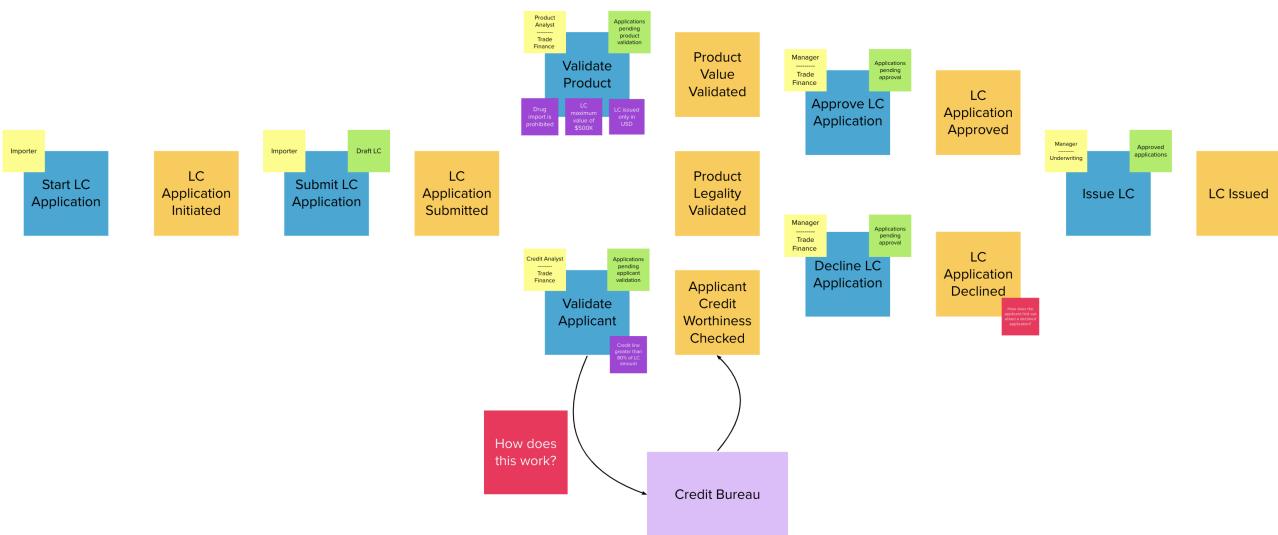
```

<dependency>
    <groupId>org.axonframework</groupId>
    <artifactId>axon-test</artifactId>
    <scope>test</scope>
    <version>${axon-framework.version}</version>
</dependency>

```

With the above set up, you should be able to run the application and start implementing the LC issuance functionality.

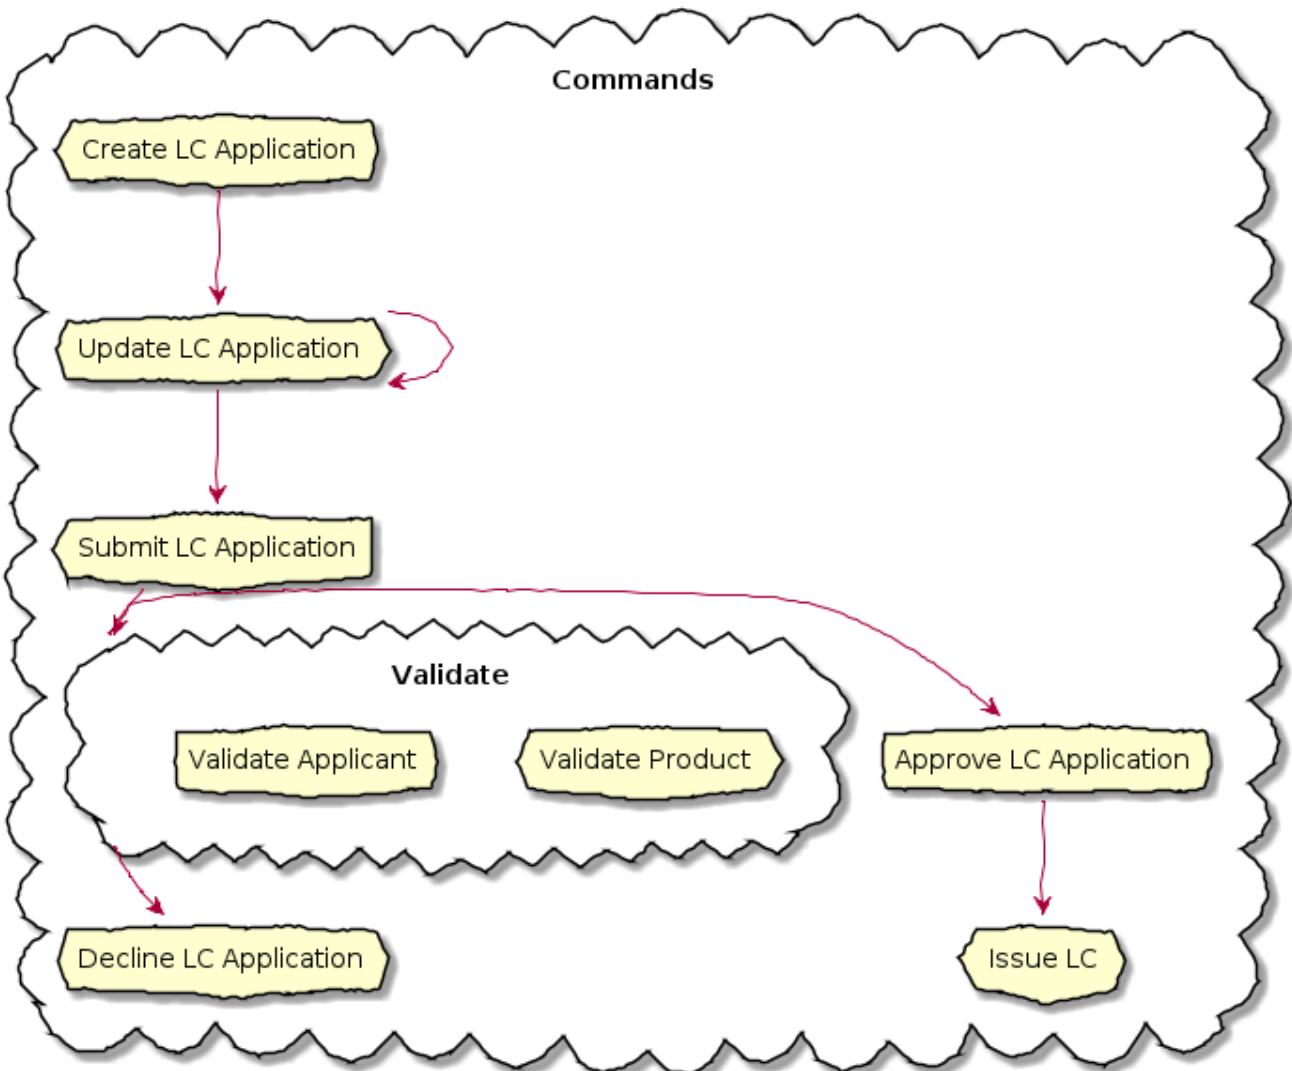
As a reminder, this is the output produced from our eventstorming session:



The blue stickies in this diagram represent commands. Let's look at how to implement these commands using the Axon framework.

5.4. Identifying Commands

From the eventstorming session, we have the following commands to start with:



Commands are always directed to an aggregate for processing (handling). This means that we need to resolve each of these commands to be handled by an aggregate. While the sender of the command does not care which component within the system handles it, we need to decide which aggregate will handle each command. It is also important to note that any given command can only be handled by a single aggregate within the system. Let's look at how to group these commands and assign them to aggregates. To be able to do that, we need to identify the aggregates in the system first.

5.4.1. Identifying Aggregates

Looking at the output of the eventstorming session, one potential grouping can be as follows:

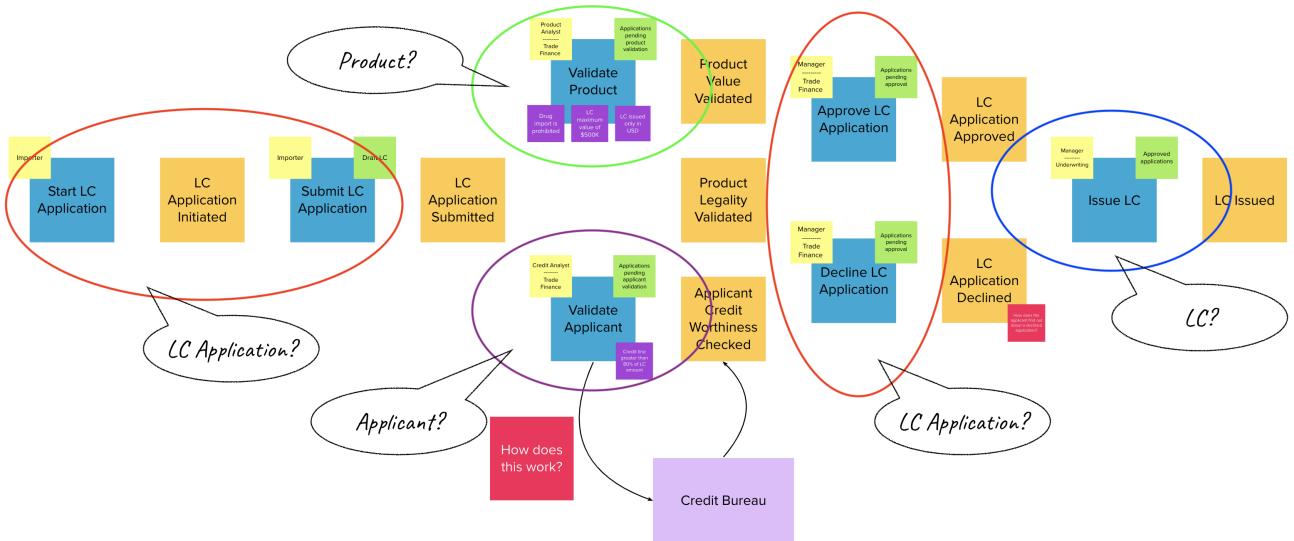


Figure 1-28. First cut attempt at aggregate design

At first glance, it appears that we have four potential entities to handle these commands:

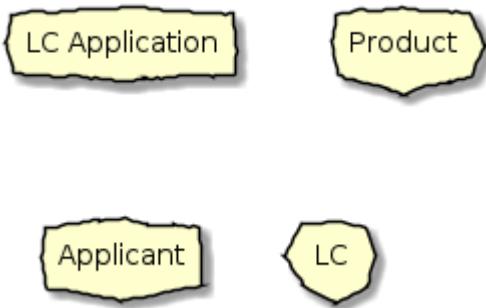


Figure 1-29. Potential aggregates at first glance

At first glance, each of these entities may be classified as aggregates in our solution. Here, the **LC Application** feels like a reasonably good choice for aggregate, given that we are building a solution to manage LC applications. However, do the others make sense to be classified as such? The **Product** and **Applicant** look like potential entities, but we need to ask ourselves if we will need to operate on these outside of the purview of the **LC Application**. If the answer is a **yes**, then **Product** and **Applicant** **may** be classified as aggregates. But both **Product** and **Applicant** do not seem to require being operated on without an enveloping **LC Application** within this bounded context. It feels that way because both product and applicant details are required to be provided as part of the LC application process. At least from what we know of the process thus far, this seems to be true. This means we are left with two potential aggregates — **LC** and **LC Application**.

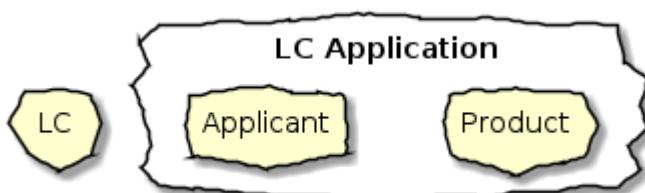


Figure 1-30. Slightly more refined aggregate structure

When we look at the output of our eventstorming session, the **LC Application** transitions to become an **LC** much later in the lifecycle. Let's work on the **LC Application** right now, and suspend further

analysis on the need for the LC aggregate to a later time.



For a more detailed explanation of the differences between aggregates, aggregate roots, entities and value objects, refer to Chapter 2.

Let's start writing our first command to see how this manifests itself in code.

5.4.2. Test-driving the system

While we have a reasonably good conceptual understanding of the system, we are still in the process of refining this understanding. Test-driving the system allows us to exercise our understanding by acting as the first client of the solution that we are producing.



This is very well illustrated in the best-selling book—*Growing Object-Oriented Software, Guided by Tests* by authors Nat Price and Steve Freeman. This is worth looking at, to gain a deeper understanding of this practice.

So let's start with the first test. To the external world, an event-driven system typically works in a manner depicted below:

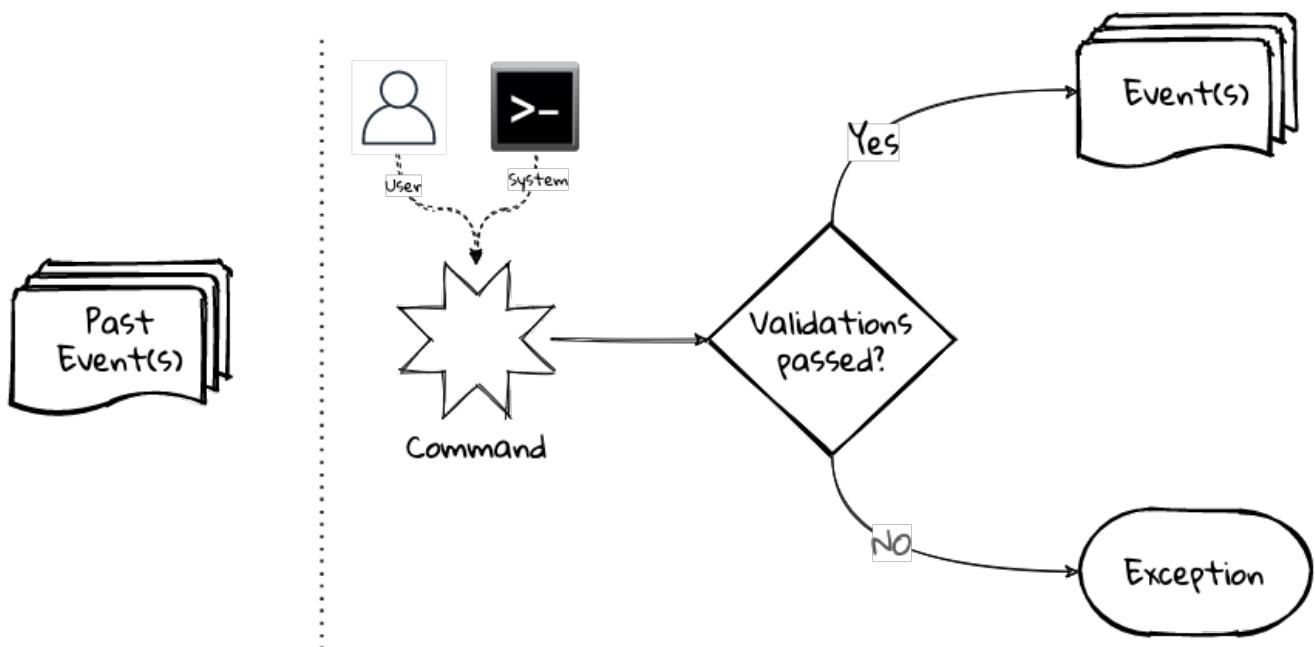


Figure 1-31. An event-driven system

An optional set of events may have occurred in the past. A command is received by the system (initiated manually by a user or automatically by a part of the system), which acts as a stimulus. The system will react in one of two ways when it handles a command:

- Emit one or more events
- Throw an exception

The Axon framework allows us to express tests in this form. This is outlined below:

```

public class LCAppliCationAggregateTests {
    private FixtureConfiguration<LCAppliCation> fixture; ①

    @BeforeEach
    void setUp() {
        fixture = new AggregateTestFixture<>(LCAppliCation.class); ②
    }

    @Test
    void shouldPublishLCAppliCationCreated() {
        fixture.given() ③

            .when(new CreateLCAppliCationCommand()) ④

            .expectEventsMatching(exactSequenceOf(
                messageWithPayload(any(LCAppliCationCreatedEvent.class)), ⑤
                andNoMore() ⑥
            ));
    } ⑦
}

```

- ① `FixtureConfiguration` is an Axon framework utility to aid testing of aggregate behaviour using a BDD style given-when-then syntax.
- ② `AggregateTestFixture` is a concrete implementation of `FixtureConfiguration` where you need to register your aggregate class—`LCAppliCation` in our case as the candidate to handle commands directed to our solution.
- ③ Since this is the start of the business process, there are no events that have occurred thus far. This is signified by the fact that we do not pass any arguments to the `given` method. In other examples we will discuss later, there will likely be events that have already occurred prior to receiving this command.
- ④ This is where we instantiate a new instance of the command object. Command objects are usually similar to data transfer objects, carrying a set of information. This command will be routed to our aggregate for handling. We will take a look at how this works in detail shortly.
- ⑤ Here we are declaring that we expect events matching an exact sequence.
- ⑥ Here we are expecting an event of type `LCAppliCationCreated` to be emitted as a result of successfully handling the command.
- ⑦ We are finally saying that we do not expect any more events—which means that we expect exactly one event to be emitted.

5.4.3. Implementing the command

The `CreateLCAppliCationCommand` in the previous simplistic example does not carry any state. Realistically, the command will likely look something like what is depicted as follows:

```

import lombok.Data;

@Data
public class CreateLCApplicationCommand { ①

    private LCApplicationId id;          ②
    private ClientId clientId;
    private Party applicant;            ③
    private Party beneficiary;
    private AdvisingBank advisingBank;   ③
    private LocalDate issueDate;
    private MonetaryAmount amount;
    private String merchandiseDescription;

}

```

- ① The command class. When naming commands, we typically use an imperative style i.e. they usually begin with a verb denoting the action required. Note that this is a data transfer object. In other words, it is simply a bag of data attributes. Also note how it is devoid of any logic (at least at the moment).
- ② The identifier for the LC Application. We are assuming client generated identifiers in this case. The topic of using server-generated versus client-generated identifiers is out of scope for the subject of this book. You may use either depending on what is advantageous in your context. Also note that we are using a strong type for the identifier `LCApplicationId` as opposed to a primitive such as a numeric or a string value. It is also common in some cases to use UUIDs as the identifier. However, we prefer using strong types to be able to differentiate between identifier types. Notice how we are using a type `ClientId` to represent the creator of the application.
- ③ The `Party` and `AdvisingBank` types are complex types to represent those concepts in our solution. Care should be taken to consistently use names that are relevant in the problem (business) domain as opposed to using names that only make sense in the solution (technology) domain. Note the attempt to make use of the *ubiquitous language* of the domain experts in both cases. This is a practice that we should always be conscious of when naming things in the system.

It is worth noting that the `merchandiseDescription` is left as a primitive `String` type. This may feel contradictory to the commentary we present above. We will address this in the upcoming section on validations.

Now let's look at what the event we will emit as a result of successfully processing the command will look like.

5.4.4. Implementing the event

In an event-driven system, mutating system state by successfully processing a command usually results in a domain event being emitted to signal the state mutation to the rest of the system. A simplified representation of a real-world `LCApplicationCreatedEvent` is shown here:

```

import lombok.Data;

@Data
public class LCAApplicationCreatedEvent { ①

    private LCAApplicationId id;
    private ClientId clientId;
    private Party applicant;
    private Party beneficiary;
    private AdvisingBank advisingBank;
    private LocalDate issueDate;
    private MonetaryAmount amount;
    private String merchandiseDescription;

}

```

- ① The event type. When naming events, we typically use names in the past tense to denote things that have already occurred and are to be accepted unconditionally as empirical facts that cannot be changed.

You will likely notice that the structure of the event is currently identical to that of the command. While this is true in this case, it may not always be that way. The amount of information that we choose to disclose in an event is context-dependent. It is important to consult with domain experts when publishing information as part of events. One may choose to withhold certain information in the event payload. For example, consider a [ChangePasswordCommand](#) which contains the newly changed password. It might be prudent to not include the changed password in the resulting [PasswordChangedEvent](#).

We have looked at the command and the resulting event in the previous test. Let's look at how this is implemented under the hood by looking at the aggregate implementation.

5.4.5. Designing the aggregate

The aggregate is the place where commands are handled and events are emitted. The good thing about the test that we have written is that it is expressed in a manner that hides the implementation details. But let's look at the implementation to be able to appreciate how we can get our tests to pass and meet the business requirement.

```

public class LCAApplication {

    @AggregateIdentifier
    private LCAApplicationId id;                                ①

    @SuppressWarnings("unused")
    private LCAApplication() {
        // Required by the framework
    }

    @CommandHandler
    public LCAApplication(CreateLCAApplicationCommand command) { ②
        // TODO: perform validations here
        AggregateLifecycle.apply(new LCAApplicationCreatedEvent(command.getId())); ③
    }

    @EventSourcingHandler
    private void on(LCAApplicationCreatedEvent event) {          ④
        this.id = event.getId();
    }
}

```

- ① The aggregate identifier for the `LCAApplication` aggregate. All aggregates are required to declare an identifier and mark it so using the `@AggregateIdentifier` annotation provided by the framework.
- ② The method that is handling the command needs to be annotated with the `@CommandHandler` annotation. In this case, the command handler happens to be the constructor of the class given that this is the first command that can be received by this aggregate. We will see examples of subsequent commands being handled by other methods later in the chapter.
- ③ The `@CommandHandler` annotation marks a method as being a command handler. The exact command that this method can handle needs to be passed as a parameter to the method. Do note that there can only be one command handler in the **entire** system for any given command.
- ④ Here, we are emitting the `LCAApplicationCreatedEvent` using the `AggregateLifecycle` utility provided by the framework. In this very simple case, we are emitting an event unconditionally on receipt of the command. In a real-world scenario, it is conceivable that a set of validations will be performed before deciding to either emit one or more events or failing the command with an exception. We will look at more realistic examples later in the chapter.
- ⑤ The need for the `@EventSourcingHandler` and its role are likely very unclear at this time. We will explain the need for this in detail in an upcoming section of this chapter.

This was a whirlwind introduction to a simple event-driven system. We still need to understand the role of the `@EventSourcingHandler`. To understand that, we will need to appreciate how aggregate persistence works and the implications it has on our overall design.

5.5. Persisting aggregates

When working with any system of even moderate complexity, we are required to make interactions durable. That is, interactions need to outlast system restarts, crashes, etc. So the need for persistence is a given. While we should always endeavour to abstract persistence concerns from the rest of the system, our persistence technology choices can have a significant impact on the way we architect our overall solution. We have a couple of choices in terms of how we choose to persist aggregate state that are worth mentioning:

1. State stored
2. Event sourced

Let's examine each of these techniques in more detail below:

5.5.1. State stored aggregates

Saving current values of entities is by far the most popular way to persist state—thanks to the immense popularity of relational databases and object-relational mapping (ORM) tools like Hibernate. And there is good reason for this ubiquity. Until recently, a majority of enterprise systems used relational databases almost as a default to create business solutions, with ORMs arguably providing a very convenient mechanism to interact with relational databases and their object representations. For example, for our `LCAApplication`, it is conceivable that we could use a relational database with a structure that would look something like below:

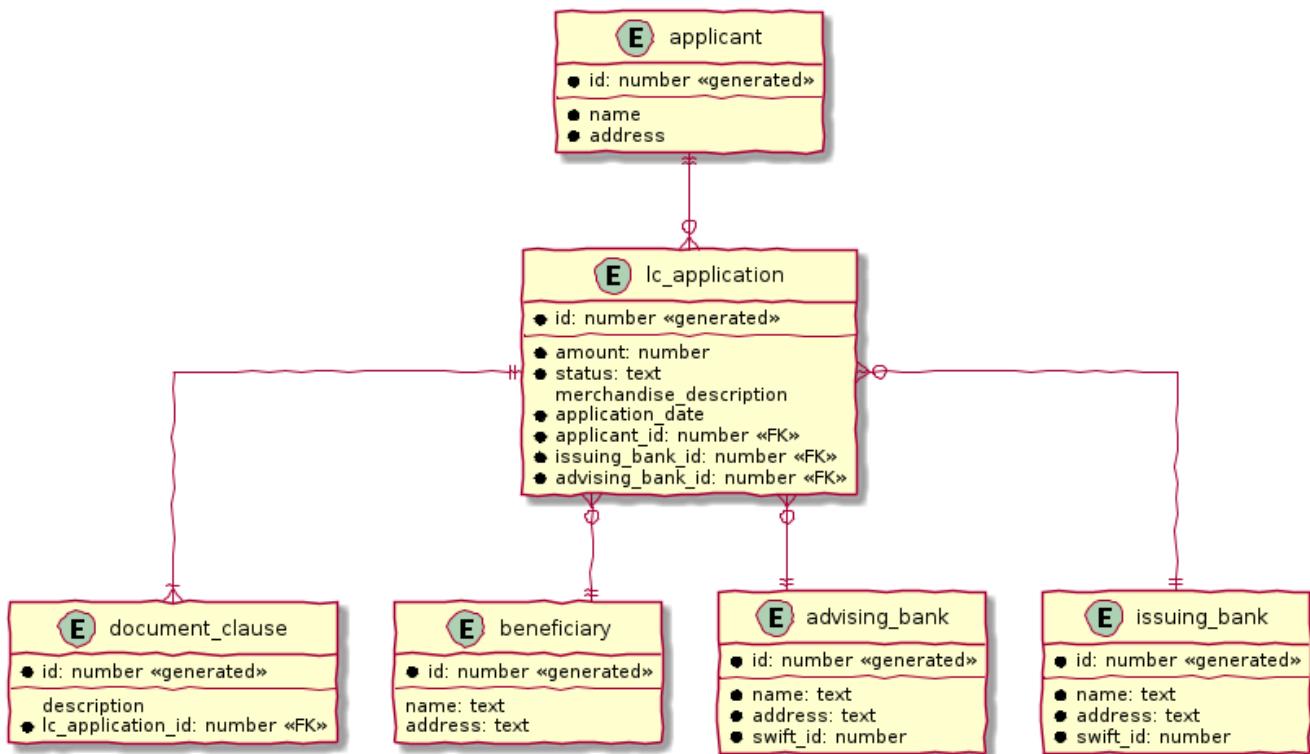


Figure 1-32. Typical entity relationship model

Irrespective of whether we choose to use a relational database or a more modern NoSQL store—for instance, a document store, key-value store, column family store, etc., the style we use to persist information remains more or less the same—which is to store the current values of the attributes

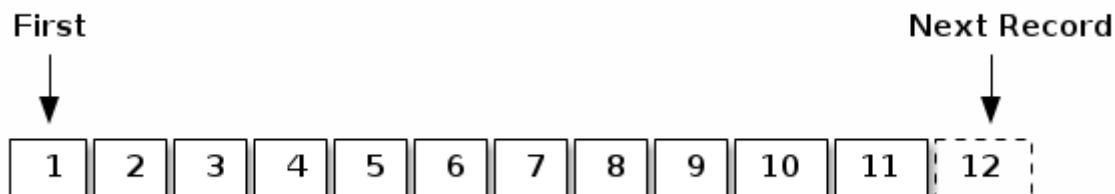
of the said aggregate/entity. When the values of attributes change, we simply overwrite old values with newer ones i.e. we store the current state of aggregates and entities—hence the name *state stored*. This technique has served us very well over the years, but there is at least one more mechanism that we can use to persist information. We will look at this in more detail below.

5.5.2. Event sourced aggregates

Developers have also been relying on logs for a variety of diagnostic purposes for a very long time. Similarly, relational databases have been employing commit logs to store information durably almost since their inception. However, developers' use of logs as a first class persistence solution for structured information in mainstream systems remains extremely rare.



A log is an extremely simple, append-only sequence of immutable records ordered by time.



Writing to a log as compared to a more complex data structure like a table is a relatively simple and fast operation and can handle extremely high volumes of data while providing predictable performance. Indeed, a modern event streaming platform like Kafka makes use of this pattern to scale to support extremely high volumes. We do feel that this can be applied to act as a persistence store when processing commands in mainstream systems because this has benefits beyond the technical advantages listed above. Consider the example of an online order flow below:

User Action	Traditional Store	Event Store
Add milk to cart	Order 123: Milk in cart	E1: Cart#123 created E2: Milk added to cart
Add white bread to cart	Order 123: Milk, White bread in cart	E1: Cart#123 created E2: Milk added to cart E3: White bread added to cart
Remove White bread from cart	Order 123: Milk in cart	E1: Cart#123 created E2: Milk added to cart E3: White bread added to cart E4: White bread removed from cart
Add Wheat bread to cart	Order 123: Milk, Wheat bread in cart	E1: Cart#123 created E2: Milk added to cart E3: White bread added to cart E4: White bread removed from cart E5: Wheat bread added to cart

User Action	Traditional Store	Event Store
Confirm cart checkout	Order 123: Ordered Milk, Wheat bread	E1: Cart#123 created E2: Milk added to cart E3: White bread added to cart E4: White bread removed from cart E5: Wheat bread added to cart E6: Order 123 confirmed

As you can see, in the event store, we continue to have full visibility of all user actions performed. This allows us to reason about these behaviors more holistically. In the traditional store, we lost the information that the user replaced white with wheat bread. While this does not impact the order itself, we lose the opportunity to gather insights from this user behavior. We recognize that this information can be captured in other ways using specialized analytical solutions, however, the event log mechanism provides a natural way to do this without requiring any additional effort, thereby reducing the complexity of the system being built. It also acts as an audit log providing full history of all events that have occurred thus far. This fits well with the essence of domain-driven design where we are constantly exploring ways in which to reduce complexity.

However, there are implications to persisting data in the form of a simple event log. Before processing any command, we will need to hydrate past events in exact order of occurrence and reconstruct aggregate state to allow us to perform validations. For example, when confirming checkout, just having the ordered set of elapsed events will not suffice. We still need to compute the exact items that are in the cart before allowing the order to be placed. This *event replay* to restore aggregate state (at least those attributes that are required to validate said command) is necessary before processing that command. For example, we need to know which items are in the cart currently before processing the [RemoveItemFromCartCommand](#). This is illustrated in the following table:

Elapsed Events	Aggregate State	Command	Event(s) Emitted
—	—	Add item: milk	E1: Cart#123 created E2: Milk added
E1: Cart#123 created E2: Milk added	Cart Items: Milk	Add item: white bread	E2: White bread added
E1: Cart#123 created E2: Milk added E3: White bread added	Cart Items: Milk, White Bread	Remove item: white bread	E3: White bread removed
E1: Cart#123 created E2: Milk added E3: White bread added E4: White bread removed	Cart Items: Milk	Add item: wheat bread	E4: Wheat bread added
E1: Cart#123 created E2: Milk added E3: White bread added E4: White bread removed E5: Wheat bread added	Cart Items: Milk Wheat bread	Confirm checkout for Cart#123	E5: Order created!

The corresponding source code for the whole scenario is illustrated in the following code snippet:

```
public class Cart {

    private boolean isNew;
    private CartItems items;
    ...

    private Cart() {                                     ①
        // Required by the framework
    }

    @CommandHandler
    public void addItem.AddItemToCartCommand command) {
        // Business validations here
        if (this.isNew) {
            apply(new CartCreatedEvent(command.getId())); ②
        }
        apply(new ItemAddedEvent(id, command.getItem()));   ②
    }

    @CommandHandler
    public void removeItem.RemoveItemFromCartCommand command) {
        // Business validations here
        apply(new ItemRemovedEvent(id, command.getItem()));
    }

    @CommandHandler
    public void checkout.ConfirmCheckoutCommand command) {
        // Business validations here
        apply(new OrderCreatedEvent(this.items));
    }

    @EventSourcingHandler
    private void on(CartCreatedEvent event) {           ③
        this.id = event.getCartId();
        this.items = new CartItems();
        this.isNew = true;
    }

    @EventSourcingHandler
    private void on(ItemAddedEvent event) {             ③
        this.items.add(event.getItem());
        this.isNew = false;
    }

    @EventSourcingHandler
    private void on(ItemRemovedEvent event) {           ③
        this.items.remove(event.getItem());
    }
}
```

```

@EventSourcingHandler
private void on(CheckoutConfirmedEvent event) {
    // ...
}

```

③

- ① Before processing any command, the aggregate loading process commences by first invoking the no-args constructor. For this reason, we need the no-args constructor to be **empty** i.e. it should **not** have any code that restores state. State restoration **must** happen only in those methods that trigger an event replay. In the case of the Axon framework, this translates to methods embellished with the `@EventSourcingHandler` annotation.
- ② It is important to note that it is possible (but not necessary) to emit **more than one event** after processing a command. This is illustrated in the first instance of the `AddItemCommand` in the previous code where we emit `CartCreatedEvent` and `ItemAddedEvent`.
- ③ The loading process continues through the invocation of event sourcing handler methods in exactly the order of occurrence for that aggregate instance.

When working with event sourced aggregates, it is very important to be disciplined about the kind of code that one can write:

Type of Method	State Restoration	Business Logic	Event Emission
<code>@CommandHandler</code>	No	Yes	Yes
<code>@EventSourcingHandler</code>	Yes	No	No

If there are a large number of events, aggregate loading can become a time-consuming operation—directly proportional to the number of elapsed events for that aggregate. There are techniques (like snapshotting) we can employ to overcome this. We will cover this in more detail in a subsequent chapter.

5.5.3. Which persistence mechanism should we choose?

Now that we have a reasonably good understanding of the two types of aggregate persistence mechanisms, it begs the question of which one we should choose. We list a few benefits of using event sourcing below:

- We get to use the events as a **natural audit log** in high compliance scenarios.
- It provides the ability to perform **more insightful analytics** on the basis of the fine-grained events data.
- It arguably produces more flexible designs when we work with a system based on **immutable events**—because the complexity of the persistence model is capped. Also, there is no need to deal with complex ORM impedance mismatch problems.
- The domain model is much more **loosely coupled** with the persistence model—enabling it to evolve mostly independently from the persistence model.
- Enables going back in time to be able to create **adhoc views and reports** without having to deal with upfront complexity.

On the flip side, these are some challenges that you might have to consider when implementing an event sourced solution:

- Event sourcing requires a **paradigm shift**. Which means that development and business teams will have to spend time and effort understanding how it works.
- The persistence model does not store state directly. This means that **adhoc querying** directly on the persistence model can be a lot more **challenging**. This can be alleviated by materializing new views, however there is added complexity in doing that.
- Event sourcing usually tends to work very well when implemented in conjunction with **CQRS** which arguably may add more complexity to the application. It also requires applications to pay closer attention to strong vs **eventual consistency** concerns.

Our experiences indicate that event sourced systems bring a lot of benefits in modern event-driven systems. However, you will need to be cognizant of the considerations presented above in the context of your own ecosystems when making persistence choices.

5.6. Policy enforcements (validations)

When processing commands, we need to enforce policies or rules. Policies come in two broad categories:

- Structural rules — those that enforce that the syntax of the dispatched command is valid.
- Domain rules — those that enforce that business rules are adhered to.

It may also be prudent to perform these validations in different layers of the system. And it is also common for some or all of these policy enforcements to be repeated in more than one layer of the system. However, the important thing to note is that before a command is successfully handled, all these policy enforcements are uniformly applied. Let's look at some examples of these in the upcoming section.

5.6.1. Structural validations

Currently, to create an LC application, one is required to dispatch a `CreateLCApplicationCommand`. While the command dictates a structure, none of it is enforced at the moment. Let's correct that.

To be able to enable validations declaratively, we will make use of the JSR-303 bean validation libraries. We can add that easily using the `spring-boot-starter-validation` dependency to our `pom.xml` file as shown here:

```
<dependency>
    <groupId>org.springframework.boot</groupId>
    <artifactId>spring-boot-starter-validation</artifactId>
</dependency>
```

Now we can add validations to the command object using the JSR-303 annotations as depicted below:

```

import lombok.Data;
import javax.validation.*;
import javax.validation.constraints.*;

@Data
public class CreateLCApplicationCommand {

    @NotNull
    private LCApplicationId id;

    @NotNull
    private ClientId clientId;

    @NotNull
    @Valid
    private Party applicant;

    @NotNull
    @Valid
    private Party beneficiary;

    @NotNull
    @Valid
    private AdvisingBank advisingBank;

    @Future
    private LocalDate issueDate;

    @Positive
    private MonetaryAmount amount;

    @NotBlank
    private String merchandiseDescription;
}

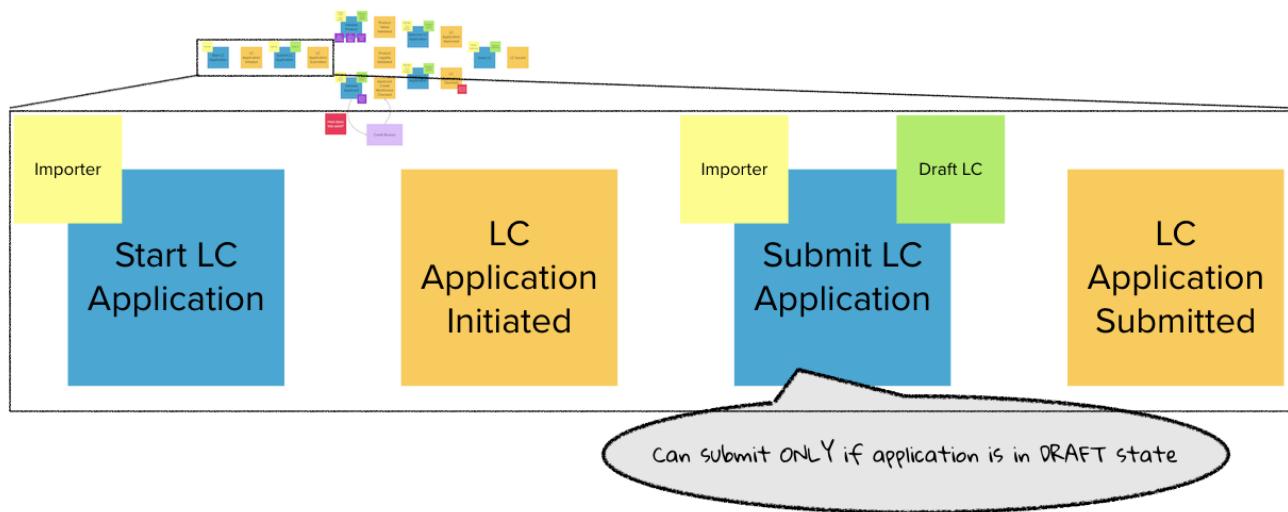
```

Most structural validations can be accomplished using the built-in validator annotations. It is also possible to create custom validators for individual fields or to validate the entire object (for example, to validate inter-dependent attributes). For more details on how to do this, please refer to the bean validation specification at <https://beanvalidation.org/2.0/> and the reference implementation at <http://hibernate.org/validator/>.

5.6.2. Business rule enforcements

Structural validations can be accomplished using information that is already available in the command. However, there is another class of validations that requires information that is not present in the incoming command itself. This kind of information can be present in one of two places: within the aggregate that we are operating on or outside of the aggregate itself, but made available within the bounded context.

Let's look at an example of a validation that requires state present within the aggregate. Consider the example of submitting an LC. While we can make several edits to the LC when it is in draft state, no changes can be made after it is submitted. This means that we can only submit an LC once. This act of submitting the LC is achieved by issuing the `SubmitLCAplicationCommand` as shown in the artifact from the eventstorming session:



Let's begin with a test to express our intent:

```

class LCAplicationAggregateTests {
    ...
    @Test
    void shouldAllowSubmitOnlyInDraftState() {
        final LCAplicationId applicationId = LCAplicationId.randomUUID();

        fixture.given(new LCAplicationCreatedEvent(applicationId)) ①
            .when(new SubmitLCAplicationCommand(applicationId)) ②
            .expectEvents(new LCAplicationSubmittedEvent(applicationId)); ③
    }
}

```

- ① Given that the `LCAplicationCreatedEvent` has already occurred—in other words, the LC application is already created.
- ② When we try to submit the application by issuing the `SubmitLCAplicationCommand` for the same application.
- ③ We expect the `LCAplicationSubmittedEvent` to be emitted.

The corresponding implementation will look something like:

```

class LCAplication {
    // ...
    @CommandHandler
    public void submit(SubmitLCAplicationCommand command) {
        apply(new LCAplicationSubmittedEvent(id));
    }
}

```

The implementation above allows us to submit an LC application unconditionally—more than once. However, we want to restrict users to be able to submit only once. To be able to do that, we need to remember that the LC application has already been submitted. We can do that in the `@EventSourcingHandler` of the corresponding events as shown below:

```

class LCAplication {
    // ...
    @EventSourcingHandler
    private void on(LCAplicationSubmittedEvent event) {
        this.state = State.SUBMITTED; ①
    }
}

```

- ① When the `LCAplicationSubmittedEvent` is replayed, we set the state of the `LCAplication` to `SUBMITTED`.

While we have remembered that the application has changed to be in `SUBMITTED` state, we are still not preventing more than one submit attempt. We can fix that by writing a test as shown below:

```

class LCAplicationAggregateTests {
    @Test
    void shouldNotAllowSubmitOnAnAlreadySubmittedLC() {
        final LCAplicationId applicationId = LCAplicationId.randomUUID();

        fixture.given(
            new LCAplicationCreatedEvent(applicationId), ①
            new LCAplicationSubmittedEvent(applicationId)) ②

            .when(new SubmitLCAplicationCommand(applicationId)) ③
            .expectException(AlreadySubmittedException.class) ④
            .expectNoEvents();
    }
}

```

- ① The `LCAplicationCreatedEvent` and `LCAplicationSubmittedEvent` have already happened—which means that the `LCAplication` has been submitted once.
- ② We now dispatch another `SubmitLCAplicationCommand` to the system.
- ③ We expect an `AlreadySubmittedException` to be thrown.

- ④ We also expect no events to be emitted.

The implementation of the command handler to make this work is shown below:

```
class LCAplication {  
    // ...  
    @CommandHandler  
    public void submit(SubmitLCAplicationCommand command) {  
        if (this.state != State.DRAFT) {  
            throw new AlreadySubmittedException("LC is already submitted!");  
        }  
        apply(new LCAplicationSubmittedEvent(id));  
    }  
}
```

- ① Note how we are using the state attribute from the `LCAplication` aggregate to perform the validation. If the application is not in `DRAFT` state, we fail with the `AlreadySubmittedException` domain exception.

Let's also look at an example where information needed to perform the validation is not part of either the command or the aggregate. Let's consider the scenario where country regulations prohibit transacting with a set of so called *sanctioned* countries. Changes to this list of countries may be affected by external factors. Hence it does not make sense to pass this list of sanctioned countries as part of the command payload. Neither does it make sense to maintain it as part of every single aggregate's state—given that it can change (albeit very infrequently). In such a case, we may want to consider making use of a command handler that is outside the confines of the aggregate class. Thus far, we have only seen examples of `@CommandHandler` methods within the aggregate. But the `@CommandHandler` annotation can appear on any other class external to the aggregate. However, in such a case, we need to load the aggregate ourselves. The Axon framework provides a `org.axonframework.modelling.command.Repository` interface to allow us to do that. It is important to note that this `Repository` is distinct from spring framework's interface that is part of the spring data libraries. An example of how this works is shown below:

```

import org.axonframework.modelling.command.Repository;

class MyCustomCommandHandler {

    private final Repository<LCApplication> repository; ①

    MyCustomCommandHandler(Repository<LCApplication> repository) {
        this.repository = repository; ①
    }

    @CommandHandler
    public void handle(SomeCommand command) {
        Aggregate<LCApplication> application = repository.load(command.getAggregateId
()); ②
        // Command handling code
    }

    @CommandHandler
    public void handle(AnotherCommand command) {
        Aggregate<LCApplication> application = repository.load(command.getAggregateId
());
        // Command handling code
    }
}

```

- ① We are injecting the Axon **Repository** to allow us to load aggregates. This was not required previously because the **@CommandHandler** annotation appeared on aggregate methods directly.
- ② We are using the **Repository** to load aggregates and work with them. The **Repository** interface supports other convenience methods to work with aggregates. Please refer to the Axon framework documentation for more usage examples.

Coming back to the sanctioned countries example, let's look at how we need to set up the test slightly differently:

```

public class CreateLCApplicationCommandHandlerTests {
    private FixtureConfiguration<LCApplication> fixture;

    @BeforeEach
    void setUp() {
        final Set<Country> sanctioned = Set.of(SOKOVIA);
        fixture = new AggregateTestFixture<>(LCApplication.class);           ①

        final Repository<LCApplication> repository = fixture.getRepository(); ②

        CreateLCApplicationCommandHandler handler =
            new CreateLCApplicationCommandHandler(repository, sanctioned);      ③
        fixture.registerAnnotatedCommandHandler(handler);                      ④
    }
}

```

- ① We are creating a new aggregate fixture as usual
- ② We are using the fixture to obtain an instance of the Axon **Repository**
- ③ We instantiate the custom command handler passing in the **Repository** instance. Also note how we inject the collection of sanctioned countries into the handler using simple dependency injection. In real life, this set of sanctioned countries will likely be obtained from external configuration.
- ④ We finally need to register the command handler with the fixture, so that it can route commands to this handler as well.

The tests for this look fairly straightforward:

```

class CreateLCApplicationCommandHandlerTests {
    // ...

    @BeforeEach
    void setUp() {
        final Set<Country> sanctioned = Set.of(SOKOVIA); ①
        fixture = new AggregateTestFixture<>(LCApplication.class);

        final Repository<LCApplication> repository = fixture.getRepository();

        CreateLCApplicationCommandHandler handler =
            new CreateLCApplicationCommandHandler(repository, sanctioned); ②
        fixture.registerAnnotatedCommandHandler(handler);
    }

    @Test
    void shouldFailIfBeneficiaryCountryIsSanctioned() {
        fixture.given()
            .when(new CreateLCApplicationCommand(randomId(), SOKOVIA)) ③
            .expectNoEvents()
            .expectException(CannotTradeWithSanctionedCountryException.class);
    }

    @Test
    void shouldCreateIfCountryIsNotSanctioned() {
        final LCApplicationId applicationId = randomId();
        fixture.given()
            .when(new CreateLCApplicationCommand(applicationId, WAKANDA)) ④
            .expectEvents(new LCApplicationCreatedEvent(applicationId));
    }
}

```

- ① For the purposes of the test, we mark the country `SOKOVIA` as a *sanctioned* country. In a more realistic scenario, this will likely come from some form external configuration (e.g. a lookup table or form of external configuration). However, this is appropriate for our unit test.
- ② We then inject this set of *sanctioned countries* into the command handler.
- ③ When the `LCApplication` is created for the sanctioned country, we expect no events to be emitted and furthermore, the `CannotTradeWithSanctionedCountryException` exception to be thrown.
- ④ Finally, when the beneficiary belongs to a non-sanctioned country, we emit the `LCApplicationCreatedEvent` to be emitted.

The implementation of the command handler is shown below:

```

import org.springframework.stereotype.Service;

@Service ①
public class CreateLCAccountCommandHandler {
    private final Repository<LCApplication> repository;
    private final Set<Country> sanctionedCountries;

    public CreateLCAccountCommandHandler(Repository<LCApplication> repository,
                                         Set<Country> sanctionedCountries) {
        this.repository = repository;
        this.sanctionedCountries = sanctionedCountries;
    }

    @CommandHandler
    public void handle(CreateLCAccountCommand command) {
        // Validations can be performed here as well ②
        repository.newInstance(() -> new LCApplication(command, sanctionedCountries));
    }
}

```

- ① We mark the class as a `@Service` to mark it as a component devoid of encapsulated state and enable auto-discovery when using annotation-based configuration or classpath scanning. As such, it can be used to perform any "plumbing" activities.
- ② Do note that the validation for the beneficiary's country being sanctioned could have been performed on line 18 as well. Some would argue that this would be ideal because we could avoid a potentially unnecessary invocation of the Axon `Repository` method if we did that. However, we prefer encapsulating business validations within the confines of the aggregate as much as possible — so that we don't suffer from the problem of creating an [anemic domain model](#)^[3].

Finally, the aggregate implementation along with the validation is shown here:

```

class LCApplication {
    // ...
    public LCApplication(CreateLCAccountCommand command, Set<Country> sanctioned)
    {
        if (sanctioned.contains(command.getBeneficiaryCountry())) { ①
            throw new CannotTradeWithSanctionedCountryException();
        }
        apply(new LCApplicationCreatedEvent(command.getId()));
    }
}

```

- ① The validation itself is fairly straightforward. We throw a `CannotTradeWithSanctionedCountryException` when the validation fails.

With the above examples, we looked at different ways to implement the policy enforcements encapsulated within the boundaries the aggregate.

5.7. Summary

In this chapter, we used the outputs of the eventstorming session and used it as a primary aid to create a domain model for our bounded context. We looked at how to implement this using the command query responsibility segregation (CQRS) architecture pattern. We looked at persistence options and the implications of using event sourced vs state stored aggregates. Finally, we rounded off by looking at a variety of ways in which to perform business validations. We looked at all this through a set of code examples using Spring boot and the Axon framework.

With this knowledge, we should be able to implement robust, well encapsulated, event-driven domain models. In the next chapter, we will look at implementing a user interface for these domain capabilities and examine a few options such as CRUD-based vs task-based UIs.

5.8. Questions

1. Can you examine the eventstorming session artifact from the last chapter, and identify the possible aggregates that would be required?
2. In your problem domain, can you determine the right approach for persisting aggregates? What are the reasons for choosing one approach over the other?
3. Based on your current understanding, would you apply CQRS architecture pattern in your solution? And how would you justify the choice to your team ?

5.9. Further reading

Title	Author	Location
CQRS	Martin Fowler	https://martinfowler.com/bliki/CQRS.html
Bootiful CQRS and Event Sourcing with Axon Framework	SpringDeveloper and Allard Buijze	https://www.youtube.com/watch?v=7e5euKxHhTE
The Log: What every software engineer should know about real-time data's unifying abstraction	Jay Kreps	https://engineering.linkedin.com/distributed-systems/log-what-every-software-engineer-should-know-about-real-time-datas-unifying
Event Sourcing	Martin Fowler	https://martinfowler.com/eaaDev/EventSourcing.html
Using a DDD Approach for Validating Business Rules	Fabian Lopez	https://www.infoq.com/articles/ddd-business-rules/
Anemic Domain Model	Martin Fowler	https://www.martinfowler.com/bliki/AnemicDomainModel.html

5.10. Answers

1. Refer to section [Section 5.4.1](#)
2. Refer to section [Section 5.5](#), note down the pros and cons of state stored and event sourced approach, and discuss the reasons for your choice with your teammates.
3. Refer to section [Section 5.2.2](#) to list down the advantages of the approach versus the traditional approach. Share the reasoning with your teammates.

[1] <https://www.lagomframework.com/>

[2] <https://eventuate.io/>

[3] <https://www.martinfowler.com/bliki/AnemicDomainModel.html>

Chapter 6. Implementing the User Interface — Task-based (20 pages)

To accomplish a difficult task, one must first make it easy.

— Marty Rubin

The essence of DDD is a lot about capturing the business process and user intent a lot more closely. In the previous chapter, we designed a set of APIs without paying a lot of attention to how those APIs would get consumed by its eventual users. In this chapter, we will design the GUI for the LC application using the [JavaFX^{\[4\]}](#) framework. As part of that, we will examine how this approach of designing APIs in isolation can cause an impedance mismatch between the producers and the consumers. We will examine the consequences of this *impedance mismatch* and how task-based UIs can help cope with this mismatch a lot better.

At the end of the chapter, you will learn how to employ DDD principles to help you build robust user experiences that are simple and intuitive. You will also learn why it may be prudent to design your backend APIs from the perspective of the consumer.

6.1. Technical Requirements

- JDK 1.8+ (We have used Java 16 to compile sample sources)
- JavaFX SDK 16 and SceneBuilder
- Maven 3.x
- Spring Boot 2.4.x
- mvvmFX 1.8 (<https://sialcasa.github.io/mvvmFX/>)
- JUnit 5.7.x (Included with spring boot)
- TestFX (for UI testing)
- OpenJFX Monocle (for headless UI testing)
- Project Lombok (To reduce verbosity)

6.2. Bootstrapping the UI

We will simply be building on top of the LC application we created in Chapter 5: Implementing Domain Logic. For detailed instructions, refer to the section on Bootstrapping the Application. In addition, we will need to add the following dependencies to the **dependencies** section of the Maven `pom.xml` file in the root directory of the project:

```
<dependencies>
    <!-- -->
    <dependency>
        <groupId>org.openjfx</groupId>
        <artifactId>javafx-controls</artifactId>
        <version>${javafx.version}</version>
    </dependency>
    <dependency>
        <groupId>org.openjfx</groupId>
        <artifactId>javafx-graphics</artifactId>
        <version>${javafx.version}</version>
    </dependency>
    <dependency>
        <groupId>org.openjfx</groupId>
        <artifactId>javafx-fxml</artifactId>
        <version>${javafx.version}</version>
    </dependency>
    <dependency>
        <groupId>de.saxsys</groupId>
        <artifactId>mvvmfx</artifactId>
        <version>${mvvmfx.version}</version>
    </dependency>
    <dependency>
        <groupId>de.saxsys</groupId>
        <artifactId>mvvmfx-spring-boot</artifactId>
        <version>${mvvmfx.version}</version>
    </dependency>
    <!-- -->
</dependencies>
```

To run UI tests, you will need to add the following dependencies:

```

<dependencies>
    <!-- -->
    <dependency>
        <groupId>org.testfx</groupId>
        <artifactId>testfx-junit5</artifactId>
        <scope>test</scope>
        <version>${testfx-junit5.version}</version>
    </dependency>
    <dependency>
        <groupId>org.testfx</groupId>
        <artifactId>openjfx-monocle</artifactId>
        <version>${openjfx-monocle.version}</version>
    </dependency>
    <dependency>
        <groupId>de.saxsys</groupId>
        <artifactId>mvvmfx-testing-utils</artifactId>
        <version>${mvvmfx.version}</version>
        <scope>test</scope>
    </dependency>
    <!-- -->
</dependencies>

```

To be able to run the application from the command line, you will need to add the `javafx-maven-plugin` to the `plugins` section of your `pom.xml`, per the following:

```

<plugin>
    <groupId>org.openjfx</groupId>
    <artifactId>javafx-maven-plugin</artifactId>
    <version>${javafx-maven-plugin.version}</version>
    <configuration>
        <mainClass>com.premonition.lc.issuance.App</mainClass>
    </configuration>
</plugin>

```

To run the application from the command line, use:

```
mvn javafx:run
```



If you are using a JDK greater than version 1.8, the JavaFX libraries may not be bundled with the JDK itself. When running the application from your IDE, you will likely need to add the following:

```
--module-path=<path-to-javafx-sdk>/lib/ \
--add-modules=javafx.controls,javafx.graphics,javafx.fxml,javafx.media
```

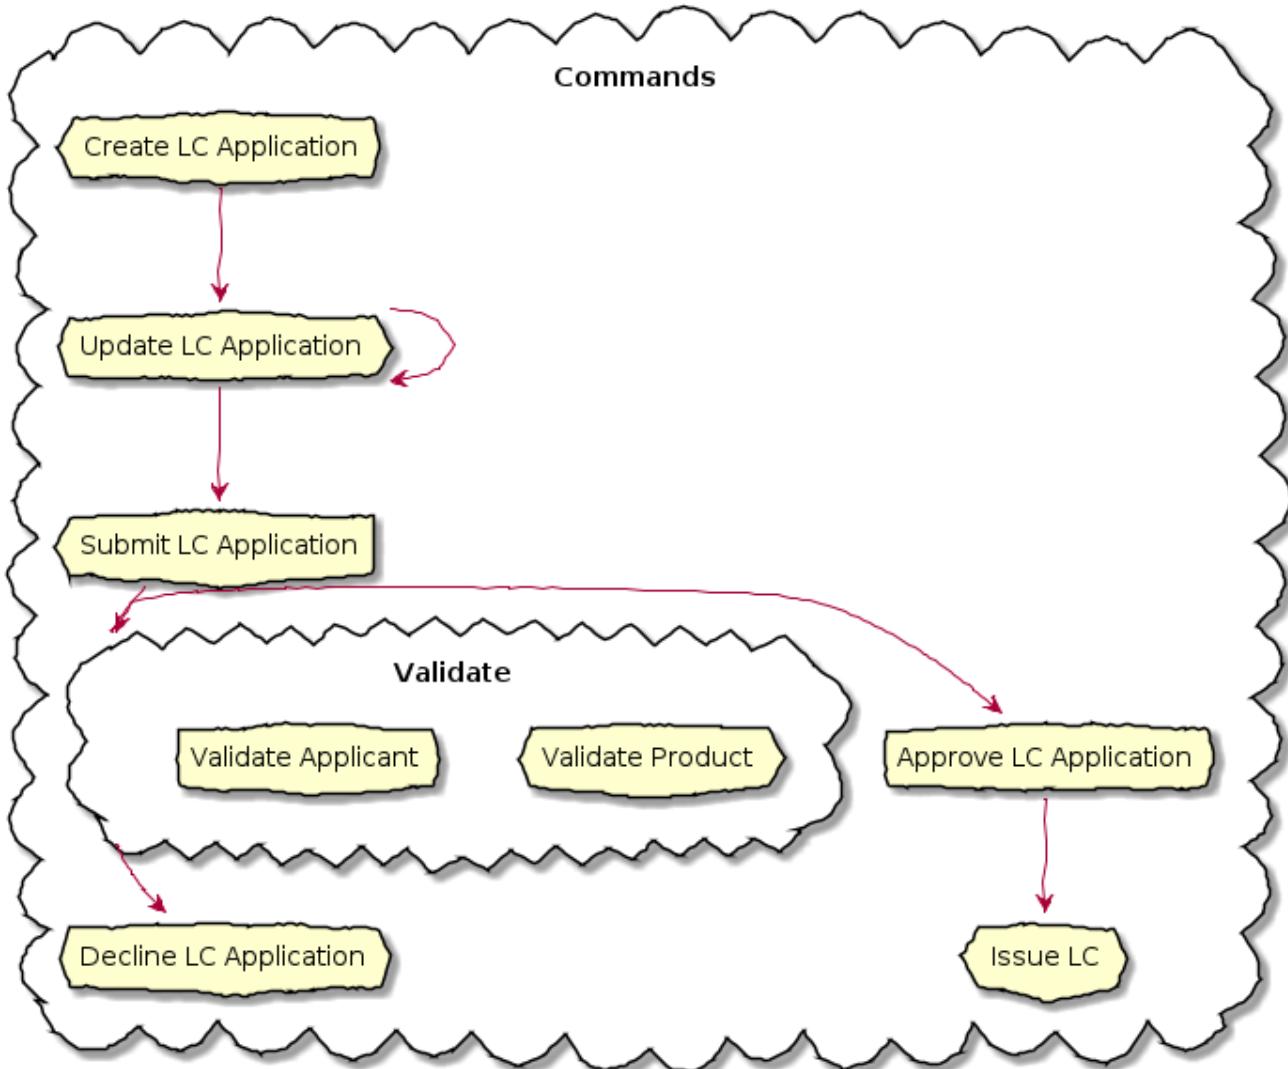


Please refer to the ch06 directory of the accompanying source code repository for the complete example.

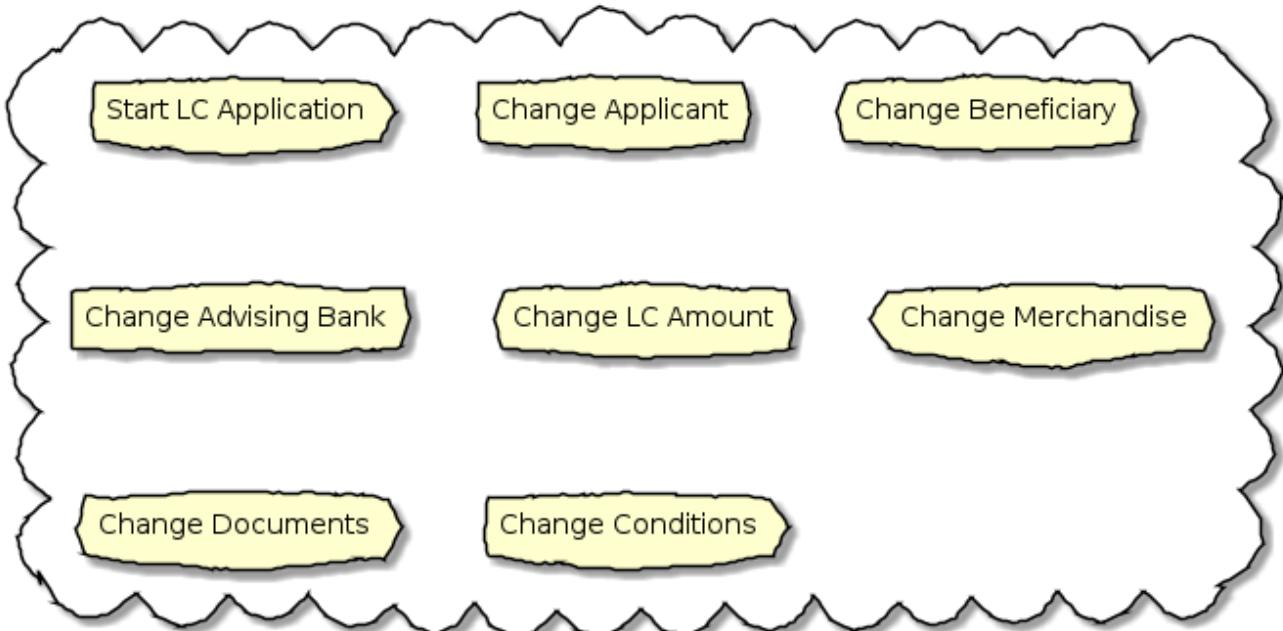
Before we dive deep into building the GUI solution, let's do a quick recap of where we left the APIs.

6.3. API Styles

If you recall from chapter 5, we created the following commands:



If you observe carefully, there seem to be commands at two levels of granularity. The "Create LC Application" and "Update LC application" are coarse grained, whereas the others are a lot more focused in terms of their intent. One possible decomposition of the coarse grained commands can be as depicted here:



In addition to just being more fine-grained than the commands in the previous iteration, the revised commands seem to better capture the user's intent. This may feel like a minor change in semantics, but can have a huge impact on the way our solution is used by its ultimate end-users. The question then is whether we should *always* prefer fine-grained APIs over coarse grained ones. The answer can be a lot more nuanced. When designing APIs and experiences, we see two main styles being employed:

- CRUD-based
- Task-based

Let's look at each of these in a bit more detail:

6.3.1. CRUD-based APIs

CRUD is an acronym used to refer to the four basic operations that can be performed on database applications: Create, Read, Update, and Delete. Many programming languages and protocols have their own equivalent of CRUD, often with slight variations in naming and intent. For example, SQL — a popular language for interacting with databases — calls the four functions Insert, Select, Update, and Delete. Similarly, the HTTP protocol has **POST**, **GET**, **PUT** and **DELETE** as verbs to represent these CRUD operations. This approach has got extended to our design of APIs as well. This has resulted in the proliferation of both CRUD-based APIs and user experiences. Take a look at the [CreateLCApplicationCommand](#) from Chapter 5:

```

import lombok.Data;

@Data
public class CreateLCAccountCommand {

    private LCAccountId id;

    private ClientId clientId;
    private Party applicant;
    private Party beneficiary;
    private AdvisingBank advisingBank;
    private LocalDate issueDate;
    private MonetaryAmount amount;
    private String merchandiseDescription;
}

```

Along similar lines, it would not be uncommon to create a corresponding `UpdateLCAccountCommand` as depicted here:

```

import lombok.Data;

@Data
public class UpdateLCAccountCommand {

    @TargetAggregateIdentifier
    private LCAccountId id;

    private ClientId clientId;
    private Party applicant;
    private Party beneficiary;
    private AdvisingBank advisingBank;
    private LocalDate issueDate;
    private MonetaryAmount amount;
    private String merchandiseDescription;
}

```

While this is very common and also very easy to grasp, it is not without problems. Here are some questions that taking this approach raises:

1. Are we allowed to change everything listed in the `update` command?
2. Assuming that is true, do they all change at the same time?
3. How do we know what exactly changed? Should we be doing a diff?
4. What if all the attributes mentioned above are not included in the `update` command?
5. What if we need to add attributes in future?
6. Is the business intent of what the user wanted to accomplish captured?

In a simple system, the answer to these questions may not matter that much. However, as system complexity increases, will this approach remain resilient to change? We feel that it merits taking a look at another approach called task-based APIs to be able to answer these questions.

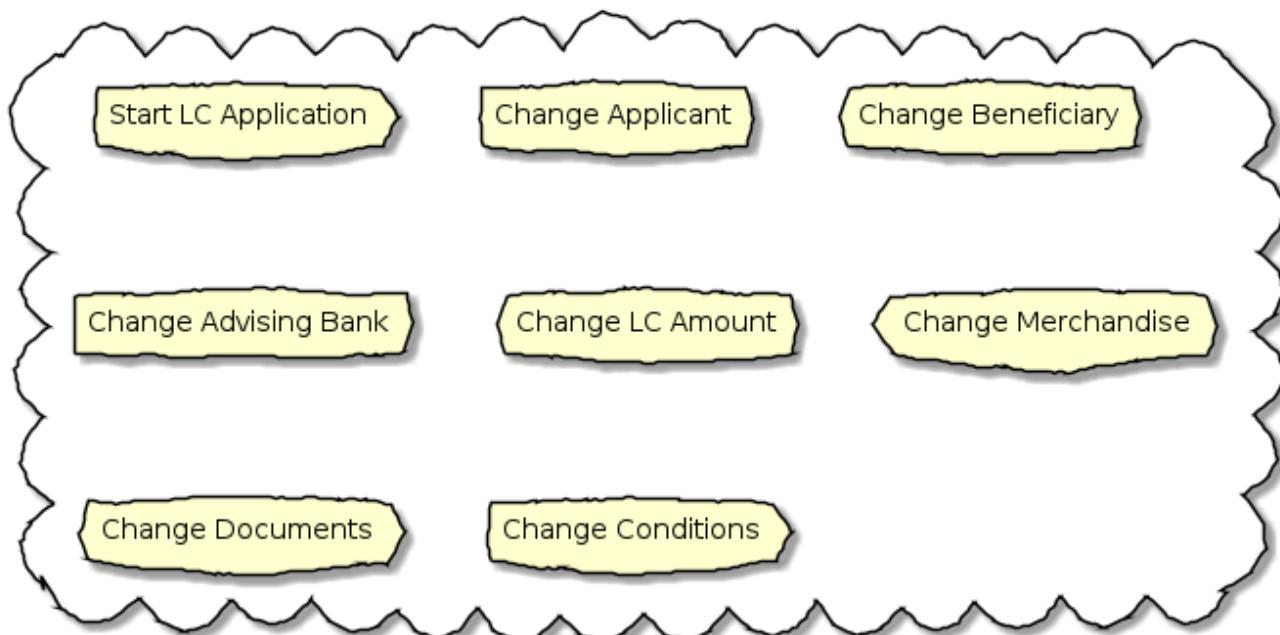
6.3.2. Task-based APIs

In a typical organization, individuals perform tasks relevant to their specialization. The bigger the organization, the higher the degree of specialization. This approach of segregating tasks according to one's specialization makes sense, because it mitigates the possibility of stepping on each others' shoes, especially when getting complex pieces of work done. For example, in the LC application process, there is a need to establish the value/legality of the product while also determining the credit worthiness of the applicant. It makes sense that each of these tasks are usually performed by individuals in unrelated departments. It also follows that these tasks can be performed independently from the other.

In terms of a business process, if we have a single `CreateLCApplicationCommand` that precedes these operations, individuals in both departments firstly have to wait for the entire application to be filled out before either can commence their work. Secondly, if either piece of information is updated through a single `UpdateLCApplicationCommand`, it is unclear what changed. This can result in a spurious notification being sent to at least one department because of this lack of clarity in the process.

Since most work happens in the form of specific tasks, it can work to our advantage if our processes and by extension, our APIs mirror these behaviors.

Keeping this in mind, let's re-examine our revised APIs for the LC application process:



While it may have appeared previously that we have simply converted our coarse-grained APIs to become more fine-grained, this in reality is a better representation of the tasks that the user intended to perform. So, in essence, task-based APIs are the decomposition of work in a manner that aligns more closely to the users' intents. With our new APIs, product validation can commence as soon as `ChangeMerchandise` happens. Also, it is unambiguously clear what the user did and what

needs to happen in reaction to the user's action.

6.3.3. Impact of API style on user experiences

6.3.4. Task-based or CRUD-based?

Task-based UIs (in contrast to CRUD-based UIs) treat user intents as first class citizens and perpetuate the spirit of DDD's ubiquitous language very elegantly.

6.4. Design pattern considerations

6.4.1. Usability

6.4.2. Testability

6.5. Implementing the save LC functionality

6.6. Invoking backend services

6.7. Summary

6.8. Questions

6.9. Further reading

Title	Author	Location
TODO	TODO	TODO

[4] <https://openjfx.com/>

Chapter 7. Implementing Queries and Projections (10 pages)

The best view comes after the hardest climb.

— Anonymous

In the section on [Section 3.8](#), we described how DDD and CQRS complement each other and how the query side (read models) can be used to create one or more representations of the underlying data. In this chapter, we will dive deeper into how we can construct read optimized representations of the data by listening to domain events. We will also look at persistence options for these read models and exposing these in the form of an HTTP API.

7.1. Consuming events

7.2. Persisting Distinct Query Models

7.3. Exposing a REST-based API for Queries

7.4. Creating Additional Views

Chapter 8. Long-Running User Flows (10 pages)

In the long run, the pessimist may be proven right, but the optimist has a better time on the trip.

— Daniel Reardon

Not all capabilities can be implemented in the form of a simple request-response interaction, requiring the management of complex state and interactions either with external systems or human-centric operations or both. In other cases, there may be a need to perform business logic at a certain time in the future. In this chapter, we will look at implementing both long-running user operations (sagas) and deadlines. We will also look at how we can keep track of the overall flow using log aggregation and distributed tracing. We will round off by looking at when/whether to choose explicit orchestrations components of implicit choreography.

8.1. Implementing Sagas

8.2. Taking Care of Deadlines

8.3. Distributed Exception Handling

8.4. Keeping Track of the Overall Flow

8.5. Deciding between Orchestration and Choreography

Chapter 9. Integrating with External Systems (15 pages)

Wholeness is not achieved by cutting off a portion of one's being, but by integration of the contraries.

— Carl Jung

Thus far, we have used DDD to implement a robust core for our application. However, most bounded contexts usually have both upstream and downstream dependencies which usually change at a pace which is different from these core components. To maintain both agility and reliability and enable loose coupling, it is important to create what DDD calls the anti-corruption layer in order to shield the core from everything that surrounds it. In this chapter, we will look at integrating with a legacy Inventory Management system. We will round off by looking at common patterns when integrating with legacy applications.

9.1. Technical Requirements

9.2. Implementing consumer-driven contracts

9.3. Exposing a REST-based API

Currently, we have a set of commands and an ability to handle them. However, there isn't an entry point to invoke these commands externally. Let's expose these via a RESTful interface.

9.4. Exposing an events-based API

9.5. Implementing an Anti-Corruption Layer

9.6. Legacy Application Migration Patterns

Part 3: Advanced Patterns

In Part 3, we will extend the application we built in Part 2 to utilize more modern, cloud native technologies. We will look at implementing an ecosystem of microservices and further extend these to be expressed to employ a serverless architecture.

Chapter 10. Distributing into Microservices (15 pages)

A distributed system is one in which the failure of a computer you didn't even know existed can render your own computer unusable.

— Leslie Lamport

We now have a working application which is bundled as a single package. In this chapter, we will distribute the UI, the command side, the query side and the saga components into distinct components. We will also look at how to test the system as a whole using the excellent [testcontainers](#) library.

10.1. Right Sizing Components

10.2. Maintaining Autonomy

10.3. Understanding the Costs of Distribution

10.4. Testing the Overall System

Chapter 11. Non-Functional Requirements (25 pages)

Sometimes I feel like I am being forgotten.

— Anonymous

While the core of the system may be met adequately, it is just as important to place focus on the operational characteristics of the system. In this chapter, we will look at common pitfalls and how to get past them.

11.1. Dealing With Eventual Consistency

11.2. Scaling the Event Store with Snapshots

11.3. Event Versioning and Upcasting

11.4. Monitoring, Metrics and Tracing

11.5. Enhancing Performance

Chapter 12. Migrating to Serverless (15 pages)

In this chapter, we will migrate the components developed thus far to adopt a serverless style of architecture.

12.1. Serverless Primer

12.2. Services as Functions

12.3. Serverless Persistence

12.4. Next Steps